

操作系统

第4章 线程 Threads

孙承杰
哈工大计算学部

E-mail: sunchengjie@hit.edu.cn
2025年秋季学期

Learning Objectives

- Understand the distinction between process and thread
- Describe the basic design issues for threads
- Explain the difference between user-level threads and kernel-level threads

Outline

❑ Processes and Threads

- ❑ Multithreading
- ❑ Thread Functionality

❑ Types of Threads

- ❑ User-Level and Kernel-Level Threads
- ❑ Other Arrangements



The basic idea is that the several components in any complex system will perform **particular subfunctions** that contribute to the **overall function**.

-- *THE SCIENCES OF THE ARTIFICIAL,*
Herbert Simon

Baidu 百科

Turing Award 1975
Nobel Prize in Economics 1978
National Medal of Science 1986
von Neumann Theory Prize 1988

Processes and Threads

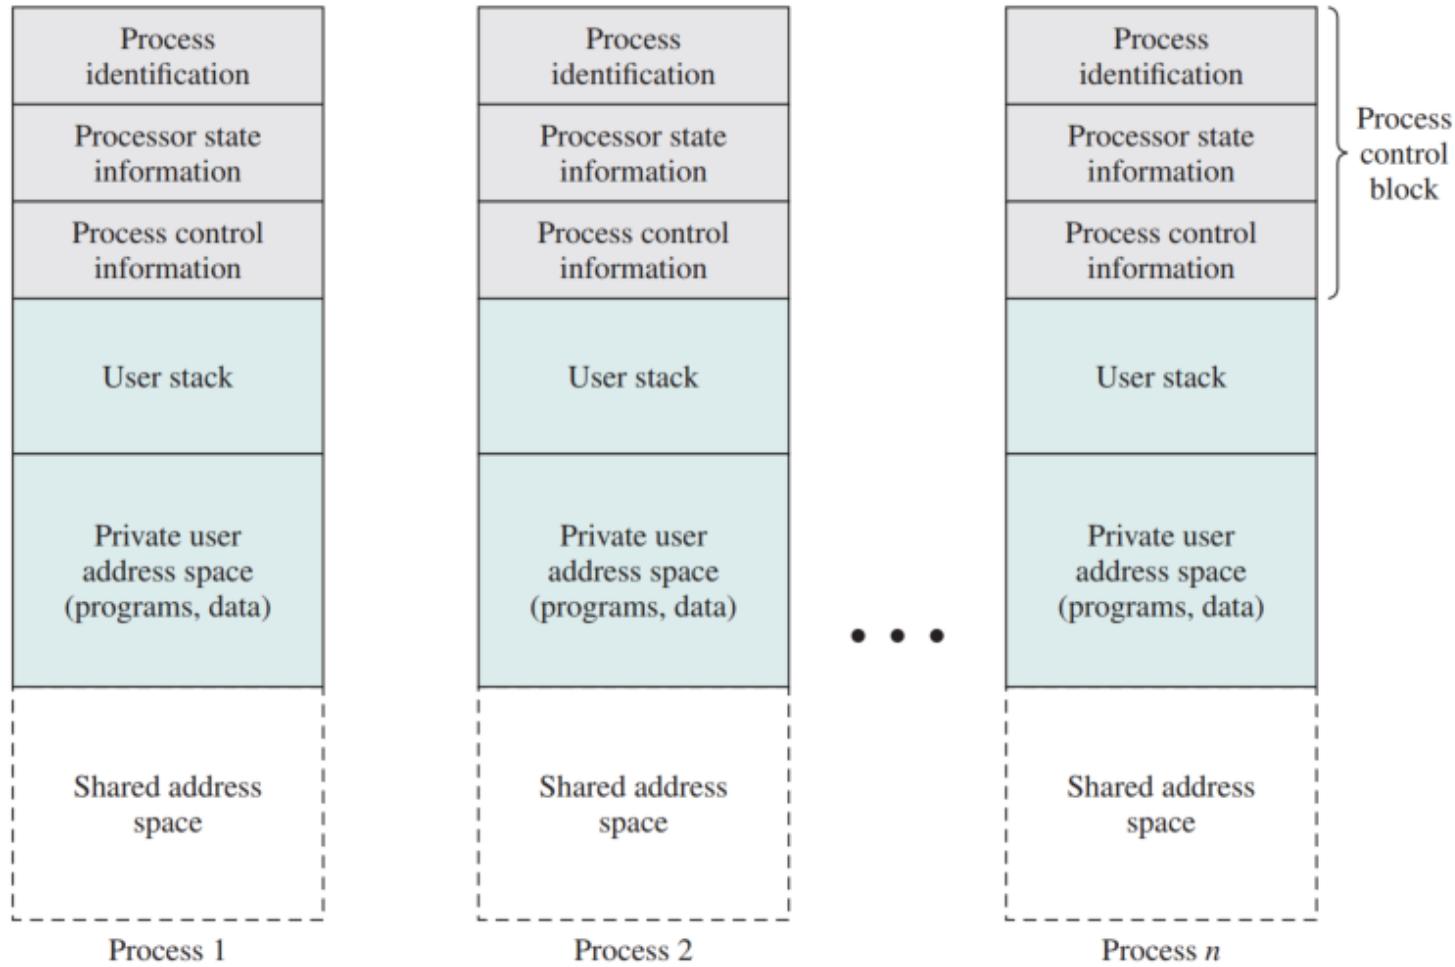
■ Resource Ownership

- a virtual address space to hold the process image
- may be allocated control or ownership of resources
 - main memory
 - I/O channels and devices
 - files
- OS performs a protection function to prevent unwanted interference between processes with respect to resources

■ Scheduling/Execution

- Follows an execution path that may be interleaved with other processes
- an execution state
 - Running, Ready, etc.
- a dispatching priority
- is the entity that is scheduled and dispatched by the OS

Structure of Process Images



User Processes in Virtual Memory

Processes and Threads

❑ Thread or lightweight process LWT

- The unit of dispatching is referred to as a thread or lightweight process

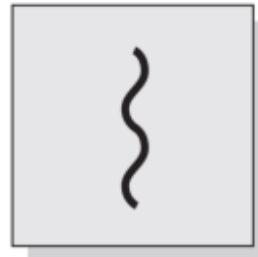
❑ Process or Task

- The unit of resource ownership is referred to as a process or task

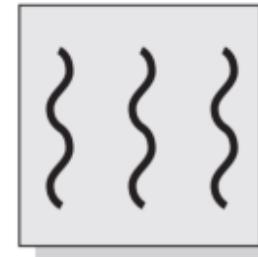
❑ Multithreading

- ❑ The ability of an OS to support multiple, concurrent paths of execution within a single process

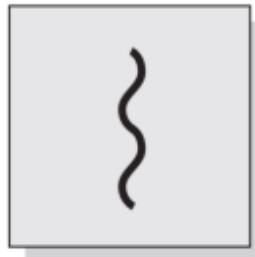
Processes and Threads



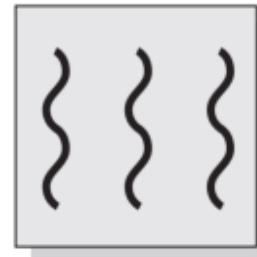
One process
One thread



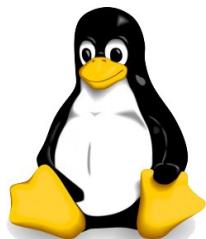
One process
Multiple threads



Multiple processes
One thread per process



Multiple processes
Multiple threads per process



= Instruction trace



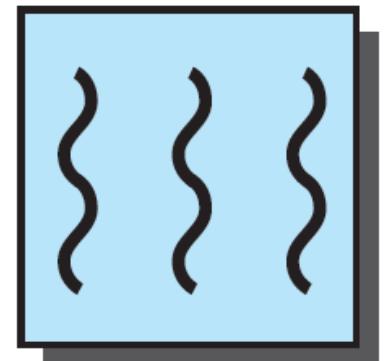
Process

□ The unit of resource allocation

- A virtual address space that holds the process image

□ The unit of protection

- access to processors
- other processes
 - for interprocess communication
- files
- I/O resources
 - devices and channels



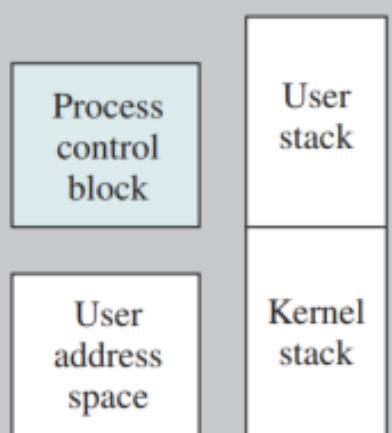
One or More Threads in a Process

□ Each thread has

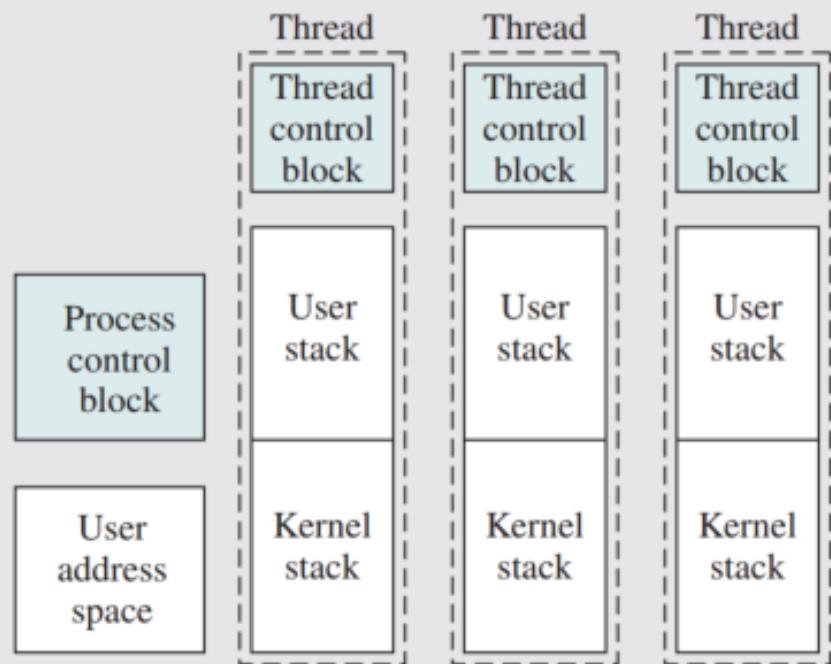
- an execution **state**
 - Running, Ready, etc.
- **saved thread context** when not running
 - one way to view a thread is **as an independent program counter** operating **within a process**
- an execution **stack**
- some **per-thread static storage** for local variables
- access to the **memory and resources** of its process
 - all threads of a process **share** this

Threads vs. Processes

Single-threaded
process model



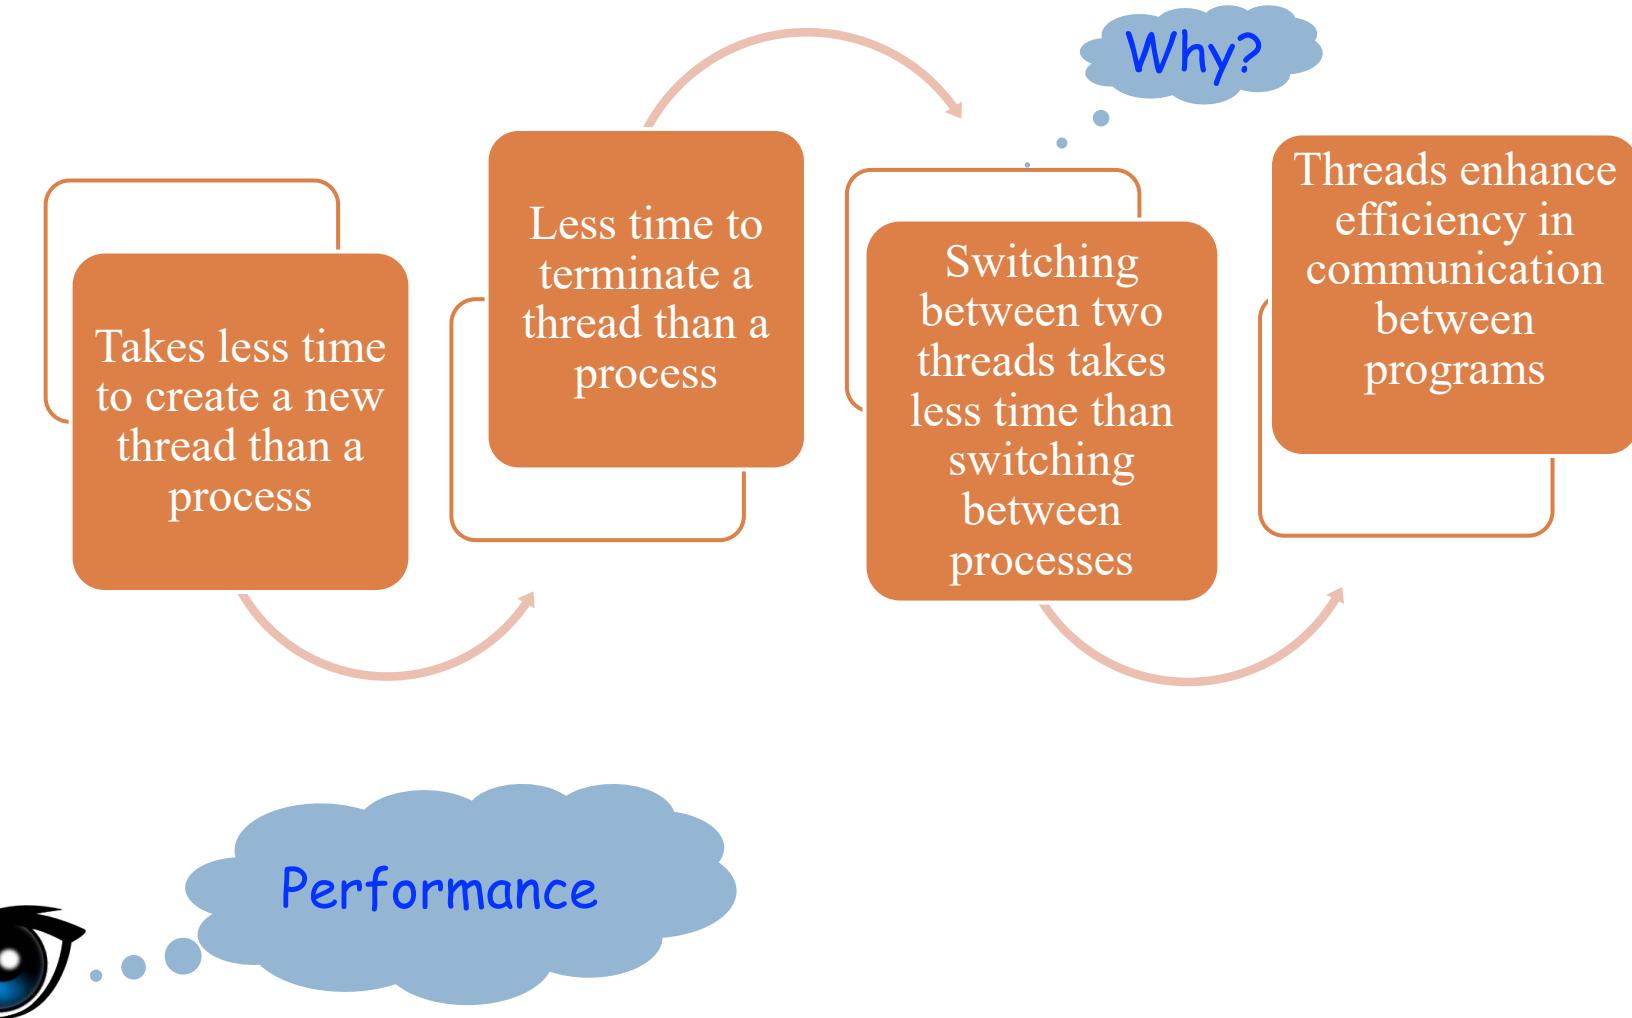
Multithreaded
process model



Process
management



Key Benefits of Threads



Thread Use in a Single-User System

- Foreground and background work
 - a spreadsheet program
- Asynchronous processing
 - as a protection against power failure
- Speed of execution
 - A multithreaded process can compute one batch of data while reading the next batch from a device
- Modular program structure
 - Programs that involve **a variety of activities or a variety of sources and destinations of input and output**

Thread Functionality

- Scheduling and dispatching is done on a thread basis
- Most of the state information dealing with execution is maintained in thread-level data structures
- Several actions that affect all of the threads in a process
 - the OS must manage at the process level
 - suspending a process involves suspending all threads of the process
 - termination of a process terminates all threads within the process

Thread Execution States

■ Key states for a thread

- Running
- Ready
- Blocked

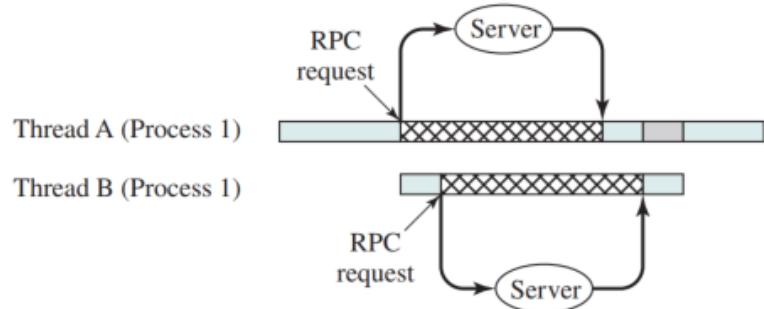
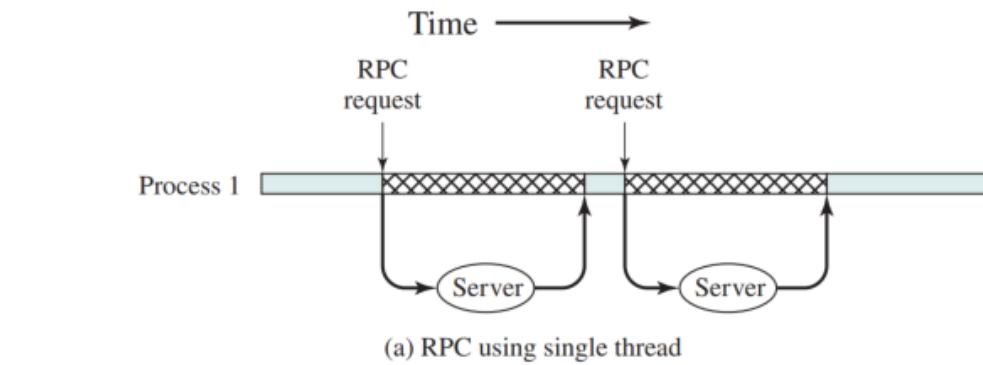
■ Thread operations associated with a change in thread state

- Spawn
- Block
- Unblock
- Finish

■ Spawn

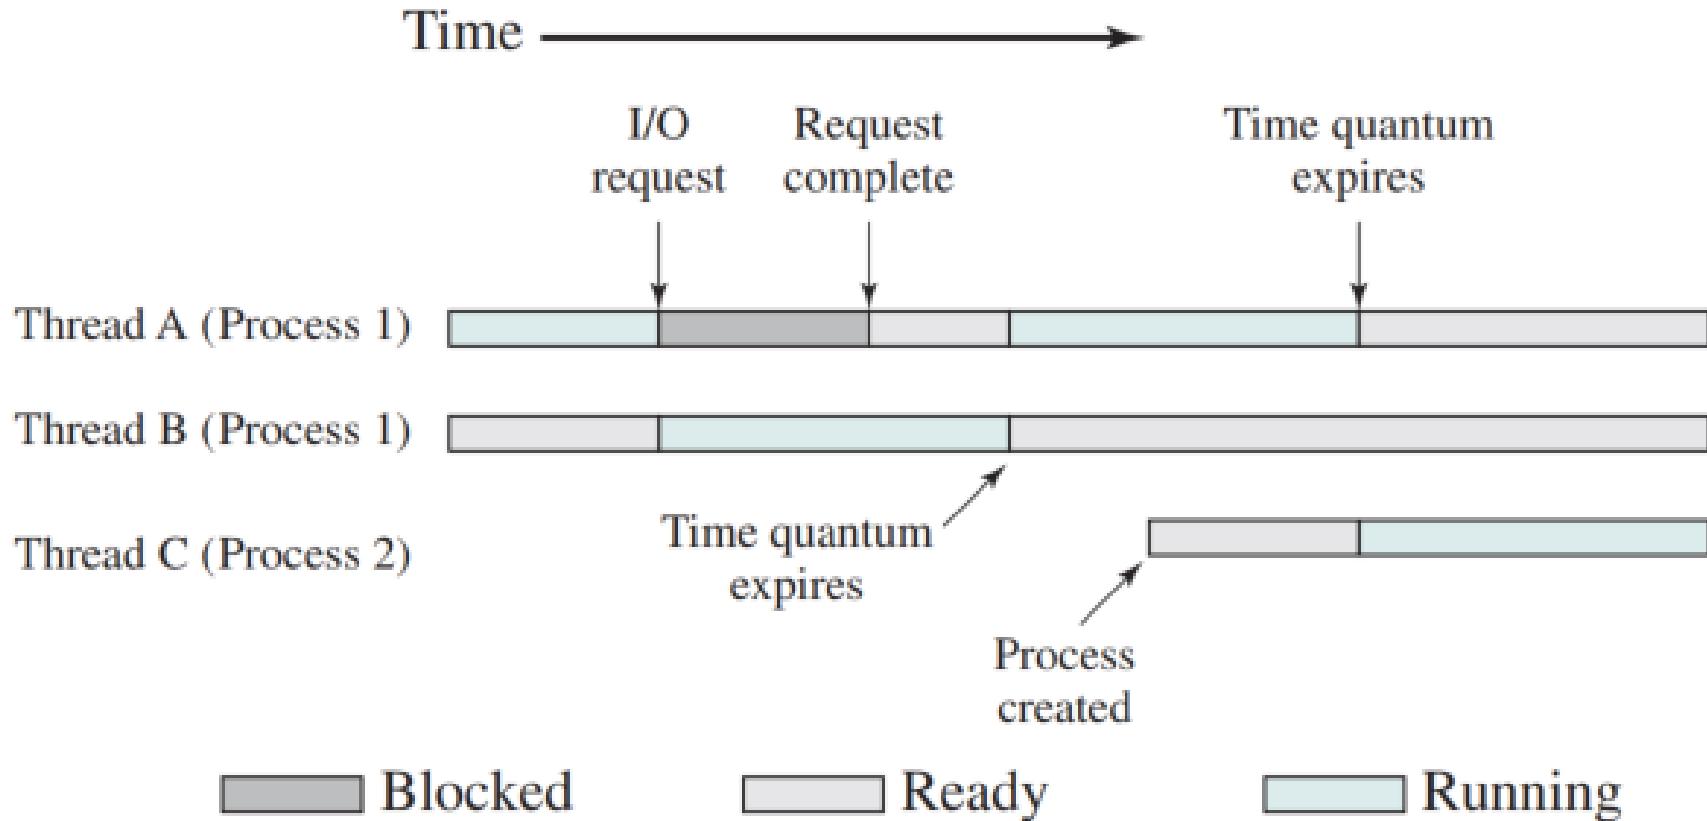
- when a new process is spawned, a thread for that process is also spawned
- a thread within process may spawn another thread within the same process
- The new thread is provided with its own register context and stack space and placed on the ready queue

Performance Benefits of Threads



- ☒ Blocked, waiting for response to RPC
- ☒ Blocked, waiting for processor, which is in use by Thread B
- ☒ Running

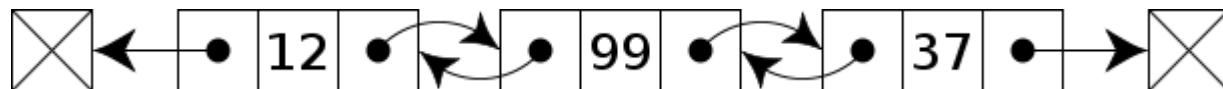
Multithreading on a Uniprocessor



Thread Synchronization

□ It is necessary to synchronize the activities of the various threads

- all threads of a process share the same address space and other resources
- any alteration of a resource by one thread affects the other threads in the same process



if two threads each try to add an element to a doubly linked list at the same time

- one element may be lost
- the list may end up malformed

Types of Threads

whether the blocking of a
thread results in the
blocking of the entire
process ?



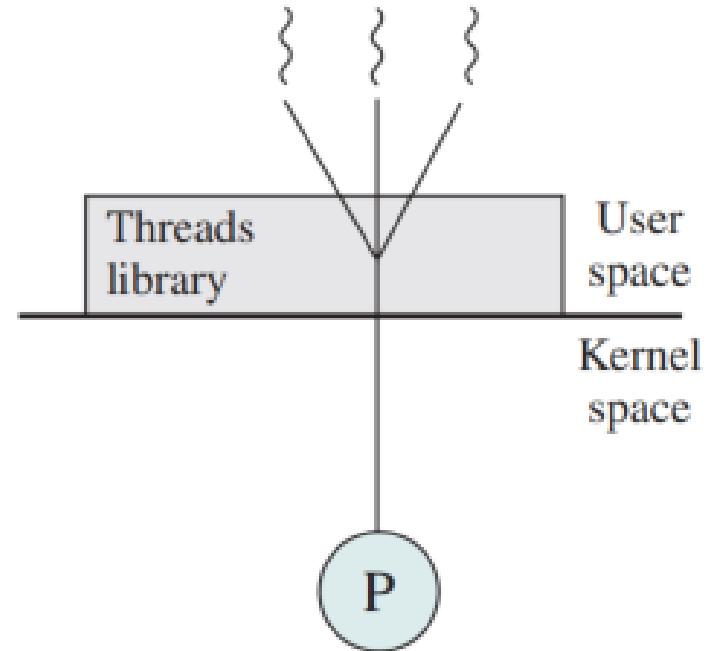
User Level Thread (ULT)
Kernel level Thread (KLT)

User-Level Threads (ULTs)

□ All thread management is done by the application

- in user space
- within a single process

□ The kernel is not aware of the existence of threads



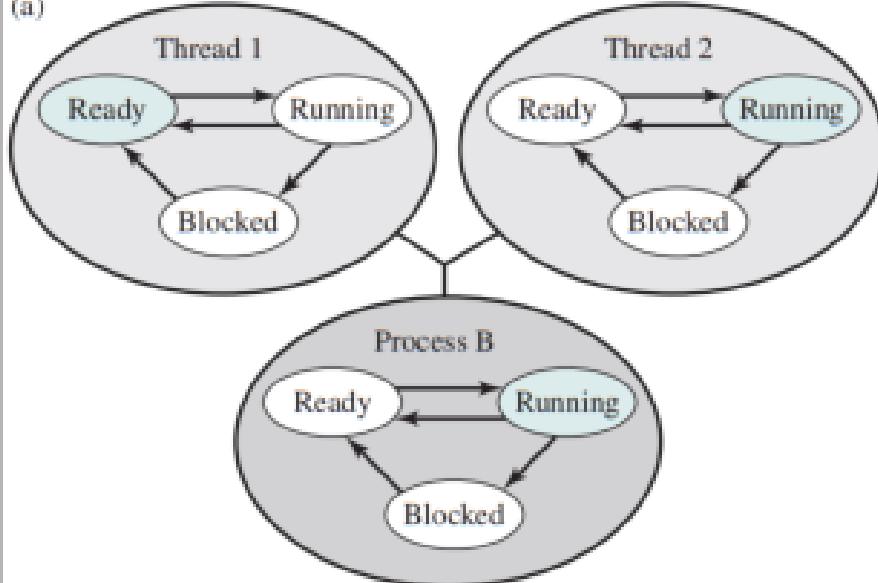
Pure user-level

{ User-level thread

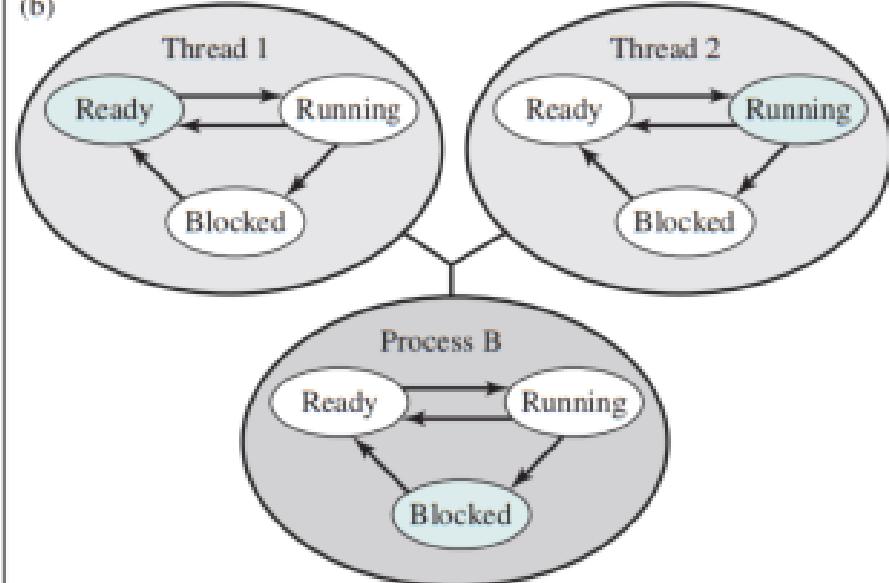
{ Kernel-level thread

P Process

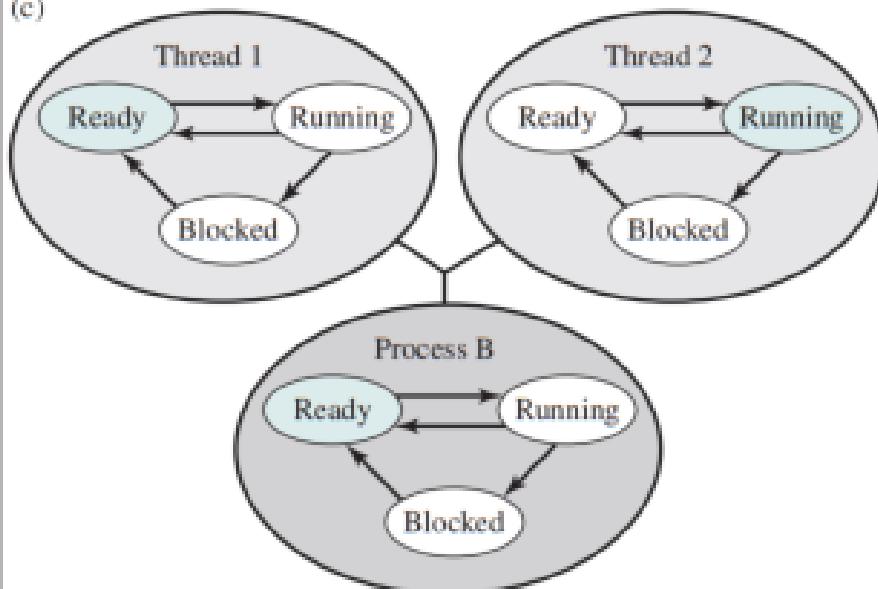
(a)



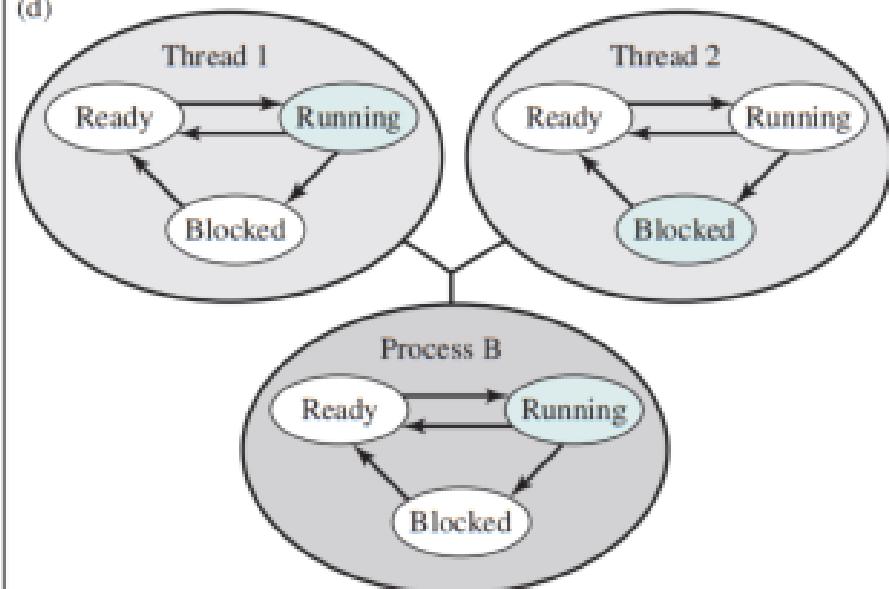
(b)



(c)



(d)



Examples of the Relationships Between User-Level Thread States and Process States

Advantages of ULTs

- Thread switching does not require kernel mode privileges
 - This saves the overhead of two mode switches
 - user to kernel; kernel back to user
- Scheduling can be application specific
 - The scheduling algorithm can be tailored to the application without disturbing the underlying OS scheduler
- ULTs can run on any OS
 - The threads library is a set of application-level functions shared by all applications

Disadvantages of ULTs

■ Disadvantages

- when a ULT executes a **system call**, not only is that thread blocked, but also all of the threads within the process are blocked
- In a pure ULT strategy, a **multithreaded application** cannot take advantage of **multiprocessing**

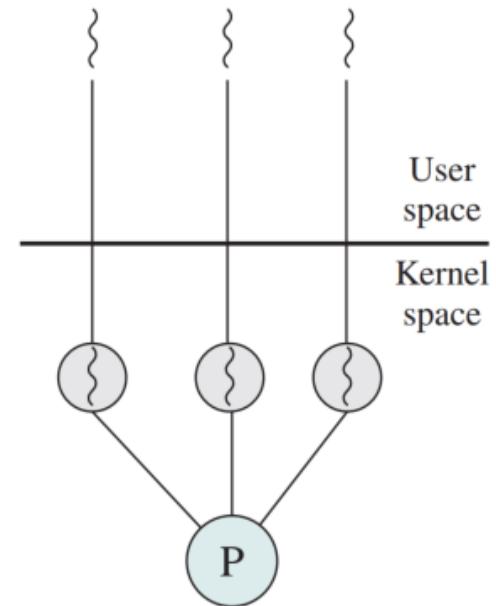
■ Overcoming

- Writing an application as **multiple processes** rather than **multiple threads**
 - Each switch becomes a **process switch** rather than a **thread switch**, resulting in much **greater overhead**
- **Jacketing**
 - converts a **blocking** system call into a **non-blocking** system call

Kernel-Level Threads (KLTs)

❑ Thread management is done by the kernel

- no thread management is done by the application
- Windows is an example of this approach



Pure kernel–level

{ User-level thread

{ Kernel-level thread

P Process

Advantages of KLTs

- The kernel can simultaneously schedule multiple threads from the same process on multiple processors
- If one thread in a process is blocked, the kernel can schedule another thread of the same process
- Kernel routines can be multithreaded

Disadvantages of KLTs

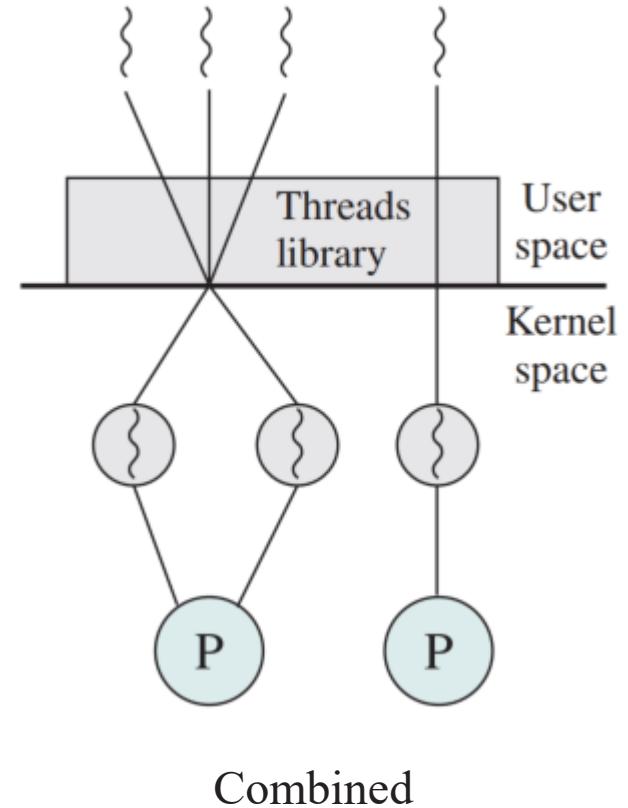
- The transfer of control from one thread to another within the same process requires a mode switch to the kernel

Thread and Process Operation Latencies (μ s)

Operation	User-Level Threads	Kernel-Level Threads	Processes
Null Fork	34	948	11,300
Signal Wait	37	441	1,840

Combined Approaches

- Thread creation is done in the user space
- Bulk of scheduling and synchronization of threads is by the application
- The multiple ULTs are mapped onto some number of KLTs
- The programmer may adjust the number of KLTs



{ User-level thread

{ Kernel-level thread

P Process

Relationship Between Threads and Processes

Threads: Processes	Description	Example Systems
1:1	Each thread of execution is a unique process with its own address space and resources.	Traditional UNIX implementations
M:1	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.	Windows NT, Solaris, Linux, OS/2, OS/390, MACH
1:M	A thread may migrate from one process environment to another. This allows a thread to be easily moved among distinct systems.	Ra (Clouds), Emerald
M:N	It combines attributes of M:1 and 1:M cases.	TRIX

Summary

■ Process

- resource ownership

■ User-level threads

- created and managed by a threads library that runs in the user space of a process
- a mode switch is not required to switch from one thread to another
- only a single user-level thread within a process can execute at a time
- if one thread blocks, the entire process is blocked

■ Thread

- program execution

■ Kernel-level threads

- threads within a process that are maintained by the kernel
- a mode switch is required to switch from one thread to another
- multiple threads within the same process can execute in parallel on a multiprocessor
- blocking of a thread does not block the entire process