

操作系统

第5章 并发性：互斥和同步 Concurrency: Mutual Exclusion and Synchronization

孙承杰
哈工大计算学部

E-mail: sunchengjie@hit.edu.cn
2025年秋季学期

Learning Objectives

- Discuss basic concepts related to concurrency
 - race conditions
 - OS concerns
 - mutual exclusion requirements
- Understand hardware approaches to supporting mutual exclusion
- Define and explain semaphores
- Define and explain monitors
- Explain the readers/writers problem

Outline

- **Mutual Exclusion: software approaches**
- **Principles of Concurrency**
 - A simple example
 - Race condition
 - Operating system concerns
 - Process interaction
 - Requirements for mutual exclusion
- **Hardware Support**
 - Interrupt disabling
- **Semaphores**
 - Mutual exclusion
 - The Producer/Consumer problem
 - Implementation of semaphores
- **Monitors**
 - Monitor with signal
- **Message Passing**
 - Synchronization
 - Addressing
 - Message format
 - Queueing discipline
 - Mutual exclusion
- **Readers/Writers Problem**
 - Readers have priority
 - Writers have priority

Designing correct routines for controlling concurrent activities proved to be one of the most difficult aspects of systems programming. The ad hoc techniques used by programmers of early multiprogramming and real-time systems were always vulnerable to subtle programming errors whose effects could be observed only when certain relatively rare sequences of actions occurred. The errors are particularly difficult to locate, since the precise conditions under which they appear are very hard to reproduce.

-- WHAT CAN BE AUTOMATED?: THE COMPUTER SCIENCE AND
ENGINEERING RESEARCH STUDY,
MIT Press, 1980

Concurrency

□ The **central themes** of OS design are all concerned with the **management of processes and threads**

➤ **Multiprogramming**

- The management of multiple processes within a **uniprocessor** system

➤ **Multiprocessing**

- The management of multiple processes within a **multiprocessor** system

➤ **Distributed processing**

- The management of multiple processes executing on **multiple, distributed** computer systems.
- clusters

Fundamental to all of these areas, and fundamental to OS design, is **concurrency**

Concurrency Arises in Three Different Contexts

□ Multiple Applications

- Multiprogramming was invented to allow processing time to be dynamically shared among a number of active applications

□ Structured Applications

- As an extension of the principles of modular design and structured programming, some applications can be effectively programmed as a set of concurrent processes

□ Operating System Structure

- OS are themselves often implemented as a set of processes or threads

Atomic operation	A function or action implemented as a sequence of one or more instructions that appears to be indivisible; that is, no other process can see an intermediate state or interrupt the operation. The sequence of instruction is guaranteed to execute as a group, or not execute at all, having no visible effect on system state. Atomicity guarantees isolation from concurrent processes.
Critical section	A section of code within a process that requires access to shared resources, and that must not be executed while another process is in a corresponding section of code.
Deadlock	A situation in which two or more processes are unable to proceed because each is waiting for one of the others to do something.
Livelock	A situation in which two or more processes continuously change their states in response to changes in the other process(es) without doing any useful work.
Mutual exclusion	The requirement that when one process is in a critical section that accesses shared resources, no other process may be in a critical section that accesses any of those shared resources.
Race condition	A situation in which multiple threads or processes read and write a shared data item, and the final result depends on the relative timing of their execution.
Starvation	A situation in which a runnable process is overlooked indefinitely by the scheduler; although it is able to proceed, it is never chosen.

Mutual Exclusion: software approaches

□ Dekker's Algorithm

- Dijkstra [DIJK65] reported an algorithm for mutual exclusion for two processes, designed by the Dutch mathematician Dekker.

```
int turn = 0;
```

```
/* PROCESS 0 */          /* PROCESS 1 *  
.  
. .  
while (turn != 0)  
    /* do nothing */ ;  
/* critical section*/;  
turn = 1;  
. .
```

(a) First attempt

Mutual Exclusion: software approaches

□ Dekker's Algorithm

```
enum boolean (false = 0; true = 1);
boolean flag[2] = 0, 0
```

```
/* PROCESS 0 *          /* PROCESS 1 *
.
.
while (flag[1])
    /* do nothing */;
flag[0] = true;
/*critical section*/;
flag[0] = false;
.

while (flag[0])
    /* do nothing */;
flag[1] = true;
/* critical section*/;
flag[1] = false;
.
```

(b) Second attempt

P0 executes the **while** statement and finds flag[1] set to false

P1 executes the **while** statement and finds flag[0] set to false

P0 sets flag[0] to true and enters its critical section

P1 sets flag[1] to true and enters its critical section

Mutual Exclusion: software approaches

□ Dekker's Algorithm

```
enum boolean (false = 0; true = 1);
boolean flag[2] = 0, 0
```

```
/* PROCESS 0 *
.
.
.
flag[0] = true;
while (flag[1])
    /* do nothing */;
    /* critical section*/;
flag[0] = false;
.

/* PROCESS 1 *
.
.
.
flag[1] = true;
while (flag[0])
    /* do nothing */;
    /* critical section*/;
flag[1] = false;
.
```

(c) Third attempt

Mutual Exclusion: software approaches

□ Dekker's Algorithm

```
enum boolean (false = 0; true = 1);  
boolean flag[2] = 0, 0
```

```
/* PROCESS 0 *          /* PROCESS 1 *  
.                .  
flag[0] = true;      flag[1] = true;  
while (flag[1]) {  while (flag[0]) {  
    flag[0] = false;  flag[1] = false;  
    /*delay */;       /*delay */;  
    flag[0] = true;   flag[1] = true;  
}  
/*critical section*/; } /* critical section*/;  
flag[0] = false;     flag[1] = false;  
.                .
```

P0 sets flag[0] to true.
P1 sets flag[1] to true.
P0 checks flag[1].
P1 checks flag[0].
P0 sets flag[0] to false.
P1 sets flag[1] to false.
P0 sets flag[0] to true.
P1 sets flag[1] to true.

(d) Fourth attempt

Dekker's Algorithm

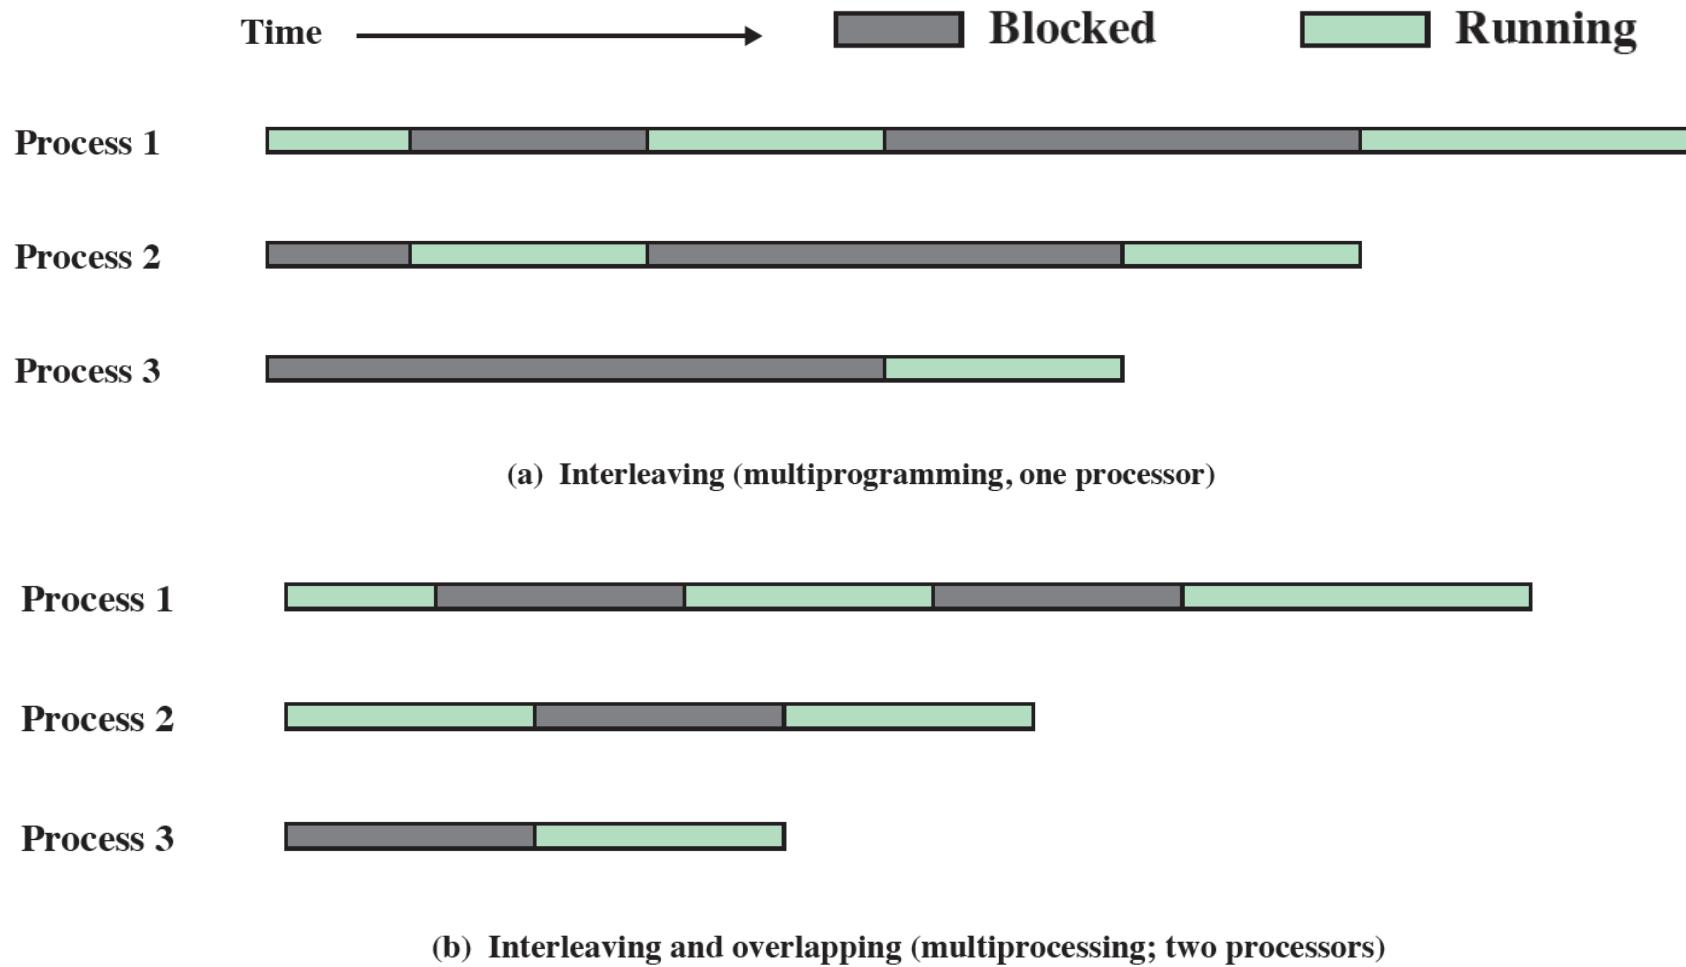
```
boolean flag [2];
int turn;
void P0()
{
    while (true) {
        flag [0] = true;
        while (flag [1]) {
            if (turn == 1)
                flag [0] = false;
            while (turn == 1) /* do nothing */;
                flag [0] = true;
        }
        /* critical section */;
        turn = 1;
        flag [0] = false;
        /* remainder */;
    }
}
void P1( )
{
    while (true) {
        flag [1] = true;
        while (flag [0]) {
            if (turn == 0) {
                flag [1] = false;
                while (turn == 0) /* do nothing */;
                    flag [1] = true;
            }
        }
        /* critical section */;
        turn = 0;
        flag [1] = false;
        /* remainder */;
    }
}
void main ()
{
    flag [0] = false;
    flag [1] = false;
    turn = 1;
    parbegin (P0, P1);
}
```

Mutual Exclusion: software approaches

□ Peterson' s Algorithm for Two Processes

```
boolean flag [2];
int turn;
void P0()
{
    while (true) {
        flag [0] = true;
        turn = 1;
        while (flag [1] && turn == 1) /* do nothing */;
        /* critical section */;
        flag [0] = false;
        /* remainder */;
    }
}
void P1()
{
    while (true) {
        flag [1] = true;
        turn = 0;
        while (flag [0] && turn == 0) /* do nothing */;
        /* critical section */;
        flag [1] = false;
        /* remainder */;
    }
}
void main()
{
    flag [0] = false;
    flag [1] = false;
    parbegin (P0, P1);
}
```

Examples of concurrent processing



Principles of Concurrency

□ Interleaving and overlapping

- can be viewed as examples of concurrent processing
- both present the same problems

□ Uniprocessor – the relative speed of execution of processes cannot be predicted

- depends on activities of other processes
- the way the OS handles interrupts
- scheduling policies of the OS

Difficulties of Concurrency

- **Sharing of global resources**
 - global variables, read/write
- **Difficult for the OS to manage the allocation of resources **optimally****
 - may lead to a **deadlock** condition
- **Difficult to **locate** programming errors**
 - results are **not deterministic and reproducible**

A simple example

```
void echo()
{
    chin = getchar();
    chout = chin;
    putchar(chout);
}
```

shared global variable



An interrupt can stop
instruction execution anywhere
in a process

How to control
access to the
shared resource?



Race Condition

- Occurs when multiple processes or threads read and write data items
- A example
 - P1 and P2, share the global variable a
 - P1 updates a to the value 1
 - P2 updates a to the value 2
 - The final result depends on the order of execution
 - the “loser” of the race is the process that updates last and will determine the final value of the variable



Operating System Concerns

- **What design and management issues are raised by the existence of concurrency?**
- **The OS must**
 - be able to keep **track** of various processes
 - allocate and de-allocate **resources** for each active process
 - Processor time, Memory, Files, I/O devices
 - **protect** the data and physical resources of each process **against** interference **by other processes**
 - **ensure** that the processes and outputs are **independent** of the **processing speed**

How to understand?

Process Interaction

Degree of Awareness	Relationship	Influence that One Process Has on the Other	Potential Control Problems
Processes unaware of each other	Competition	<ul style="list-style-type: none">• Results of one process independent of the action of others• Timing of process may be affected	<ul style="list-style-type: none">• Mutual exclusion• Deadlock (renewable resource)• Starvation
Processes indirectly aware of each other (e.g., shared object)	Cooperation by sharing	<ul style="list-style-type: none">• Results of one process may depend on information obtained from others• Timing of process may be affected	<ul style="list-style-type: none">• Mutual exclusion• Deadlock (renewable resource)• Starvation• Data coherence
Processes directly aware of each other (have communication primitives available to them)	Cooperation by communication	<ul style="list-style-type: none">• Results of one process may depend on information obtained from others• Timing of process may be affected	<ul style="list-style-type: none">• Deadlock (consumable resource)• Starvation

Resource Competition

- Concurrent processes come into conflict when they are competing for use of the same resource
 - processor time
 - memory
 - I/O devices
- In the case of competing processes three control problems must be faced
 - the need for mutual exclusion
 - deadlock
 - starvation

Mutual Exclusion

Illustration of Mutual Exclusion

```
/* PROCESS 1 * void P1
{
    while (true) {
        /* preceding code */;
        entercritical (Ra);
        /* critical section */;
        exitcritical (Ra);
        /* following code */;
    }
}
```

```
/* PROCESS 2 * void P2
{
    while (true) {
        /* preceding code */;
        entercritical (Ra);
        /* critical section */;
        exitcritical (Ra);
        /* following code */;
    }
}
```

• • •

```
/* PROCESS n * void Pn
{
    while (true) {
        /* preceding code */;
        entercritical (Ra);
        /* critical section */;
        exitcritical (Ra);
        /* following code */;
    }
}
```

entercritical (Ra);
/* critical section */;
exitcritical (Ra);

Requirements for Mutual Exclusion

- Mutual exclusion must be enforced
 - Only one process at a time is allowed into its critical section
- A process that halts must do so without interfering with other processes
- No deadlock or starvation
- A process must not be denied access to a critical section when there is no other process using it
- No assumptions are made about relative process speeds or number of processes
- A process remains inside its critical section for a finite time only

What is the solution ?



Hardware Support

■ Interrupt Disabling

- uniprocessor system
- disabling interrupts guarantees **mutual exclusion**

```
while (true) {  
    /* disable interrupts */;  
    /* critical section */;  
    /* enable interrupts */;  
    /* remainder */;  
}
```

■ Disadvantages

- the efficiency of execution could be noticeably degraded
- because the processor is limited in its ability to interleave processes
- this approach will **not work** in a multiprocessor architecture

Special Machine Instructions

■ Compare&Swap Instruction

- also called a compare and exchange instruction, can be defined as follows [HERL90]

```
int compare_and_swap (int *word, int testval, int newval)
{
    int oldval;
    oldval = *word;
    if (oldval == testval) *word = newval;
    return oldval;
}
```

Special Machine Instructions

■ Exchange Instruction

- exchanges the contents of a register with that of a memory location.

```
void exchange (int *register, int *memory)
{
    int temp;
    temp = *memory;
    *memory = *register;
    *register = temp;
}
```

Special Machine Instructions

```
/* program mutual exclusion */
const int n = /* number of processes */;
int bolt;
void P(int i)
{
    while (true) {
        while (compare_and_swap(bolt, 0, 1) == 1)
            /* do nothing */;
        /* critical section */;
        bolt = 0;
        /* remainder */;
    }
}
void main() {
    bolt = 0;
    parbegin (P(1), P(2), ..., P(n));
}
```

(a) Compare and swap instruction

```
/* program mutual exclusion */
int const n = /* number of processes */;
int bolt;
void P(int i)
{
    while (true)
        int keyi = 1;
        do exchange (&keyi, &bolt)
        while (keyi != 0);
        /* critical section */;
        bolt = 0;
        /* remainder */;
    }
}
void main() {
    bolt = 0;
    parbegin (P(1), P(2), ..., P(n));
}
```

(b) Exchange instruction

Special Machine Instructions

■ Advantages

- Applicable to any number of processes on either a single processor or multiple processors **sharing** main memory
- Simple and easy to verify
- It can be used to support **multiple critical sections**
 - each critical section can be defined by its own variable

■ Disadvantages

- **Busy-waiting** is employed
 - thus while a process is waiting for access to a critical section it continues to consume processor time
- Starvation is possible
 - when a process leaves a critical section and more than one process is waiting
- Deadlock is possible

Common Concurrency Mechanisms

Semaphore

An integer value used for signaling among processes. Only three operations may be performed on a semaphore, all of which are atomic: initialize, decrement, and increment. The decrement operation may result in the blocking of a process, and the increment operation may result in the unblocking of a process. Also known as a **counting semaphore** or a **general semaphore**.

Binary semaphore

A semaphore that takes on only the values 0 and 1.

Mutex

Similar to a binary semaphore. A key difference between the two is that the process that locks the mutex (sets the value to 0) must be the one to unlock it (sets the value to 1).

Condition variable

A data type that is used to block a process or thread until a particular condition is true.

Monitor

A programming language construct that encapsulates variables, access procedures, and initialization code within an abstract data type. The monitor's variable may only be accessed via its access procedures and only one process may be actively accessing the monitor at any one time. The access procedures are critical sections. A monitor may have a queue of processes that are waiting to access it.

Event flags

A memory word used as a synchronization mechanism. Application code may associate a different event with each bit in a flag. A thread can wait for either a single event or a combination of events by checking one or multiple bits in the corresponding flag. The thread is blocked until all of the required bits are set (AND) or until at least one of the bits is set (OR).

Mailboxes/messages

A means for two processes to exchange information and that may be used for synchronization.

Spinlocks

Mutual exclusion mechanism in which a process executes in an infinite loop waiting for the value of a lock variable to indicate availability

Semaphore

- The first major advance in dealing with the problems of concurrent processes came in 1965 with Dijkstra's treatise

The fundamental principle is this: Two or more processes can cooperate by means of simple signals, such that a process can be forced to stop at a specified place until it has received a specific signal. Any complex coordination requirement can be satisfied by the appropriate structure of signals

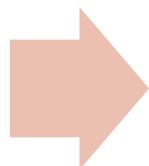


Edsger Wybe Dijkstra
Netherlands – 1972

Semaphore

- May be initialized to a **nonnegative integer** value
- The **semWait** operation **decrements** the value
- The **semSignal** operation **increments** the value

A variable that has an integer value upon which only three operations are defined:



There is no way to inspect or manipulate semaphores other than these three operations

Consequences

□ [DOWN16] points out three interesting consequences of the semaphore definition

- There is no way to know before a process decrements a semaphore whether it will **block or not**
- There is no way to know which process, if either, will **continue immediately** on a uniprocessor system
- You don't necessarily know whether another process is waiting, so the **number of unblocked processes** may be **zero or one**

A Definition of Semaphore Primitives

```
struct semaphore {  
    int count;  
    queueType queue;  
};  
void semWait(semaphore s)  
{  
    s.count--;  
    if (s.count < 0) {  
        /* place this process in s.queue */;  
        /* block this process */;  
    }  
}  
void semSignal(semaphore s)  
{  
    s.count++;  
    if (s.count<= 0) {  
        /* remove a process P from s.queue */;  
        /* place process P on ready list */;  
    }  
}
```

Binary Semaphore Primitives

```
struct binary_semaphore {  
    enum {zero, one} value;  
    queueType queue;  
};  
void semWaitB(binary_semaphore s)  
{  
    if (s.value == one)  
        s.value = zero;  
    else {  
        /* place this process in s.queue */;  
        /* block this process */;  
    }  
}  
void semSignalB(semaphore s)  
{  
    if (s.queue is empty())  
        s.value = one;  
    else {  
        /* remove a process P from s.queue */;  
        /* place process P on ready list */;  
    }  
}
```

Strong/Weak Semaphores

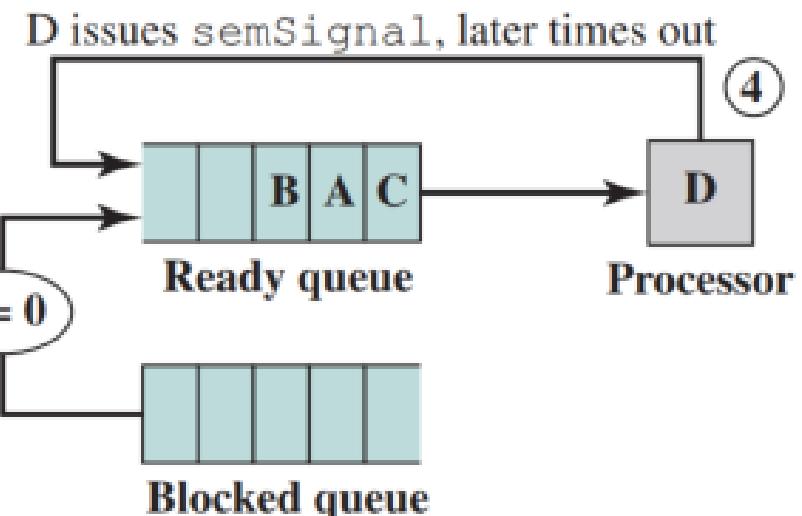
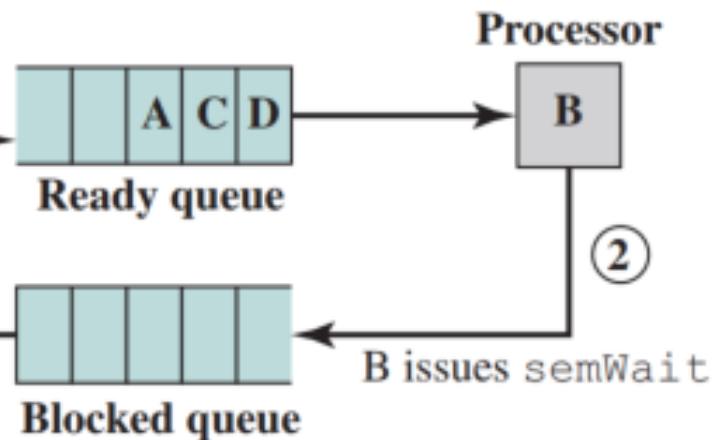
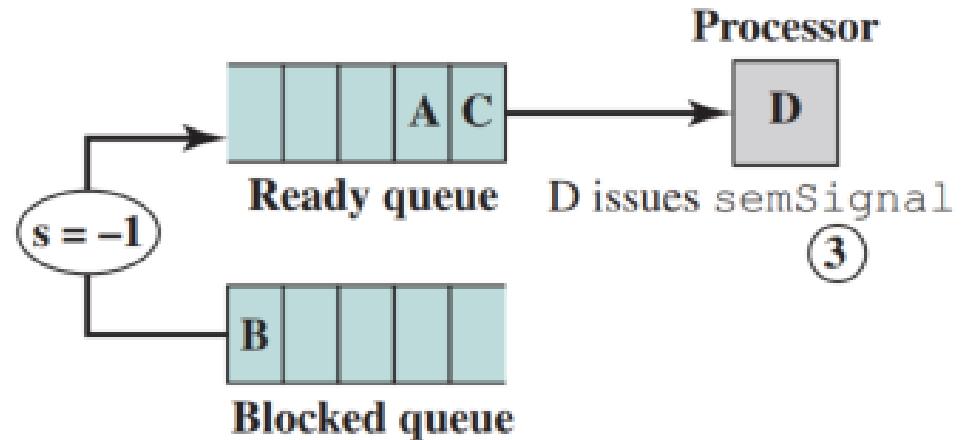
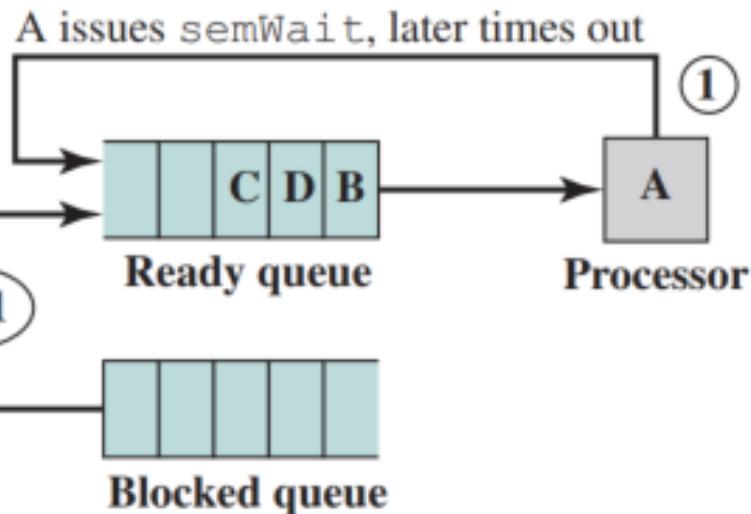
- ❑ A queue is used to hold processes waiting on the semaphore

Strong Semaphores

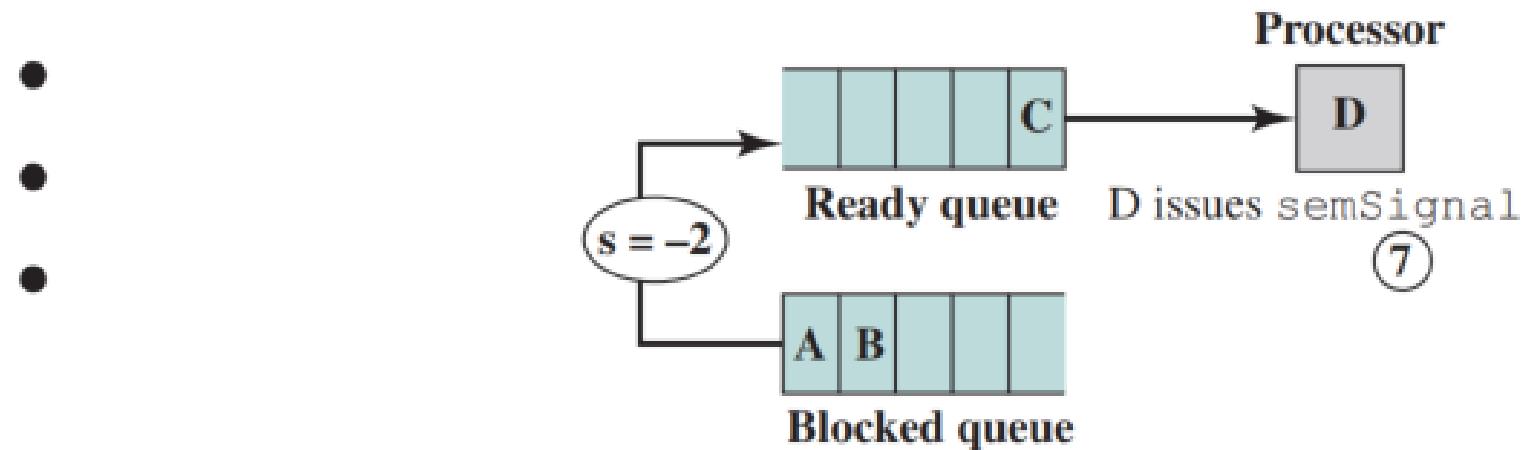
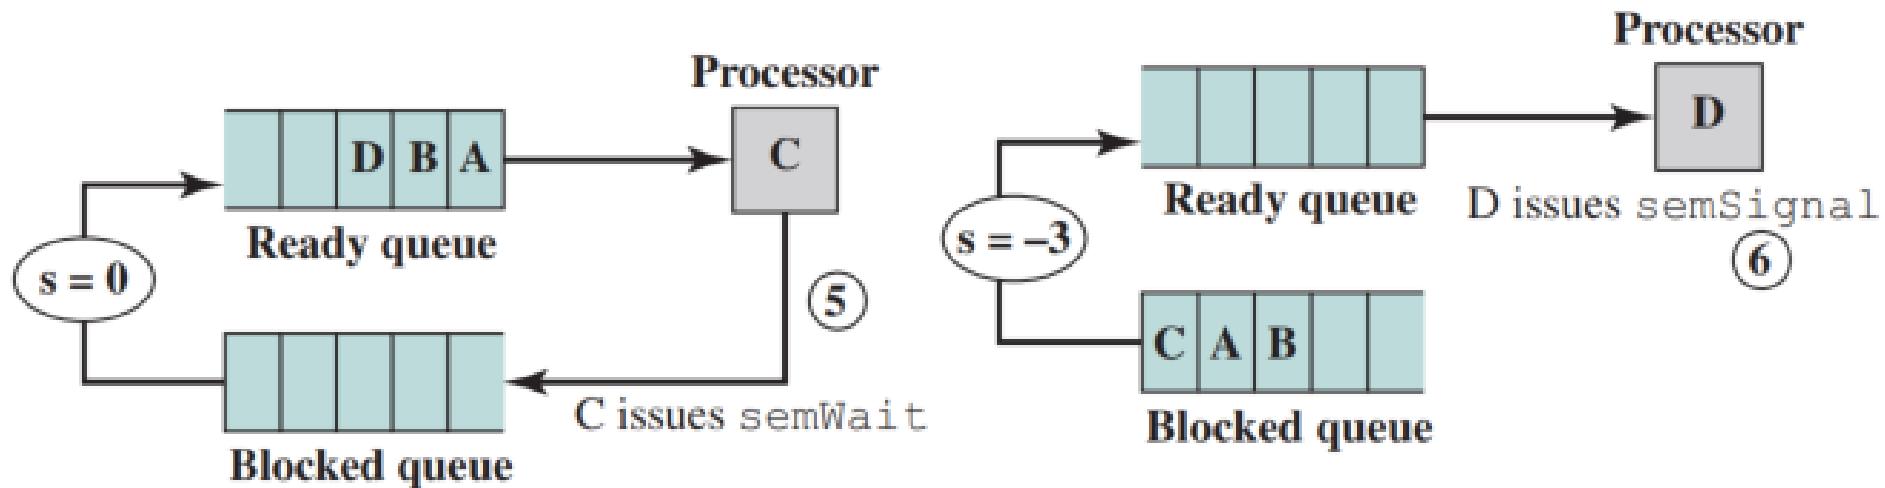
- the process that has been blocked the longest is released from the queue first (FIFO)

Weak Semaphores

- the order in which processes are removed from the queue is not specified



Example of Semaphore Mechanism

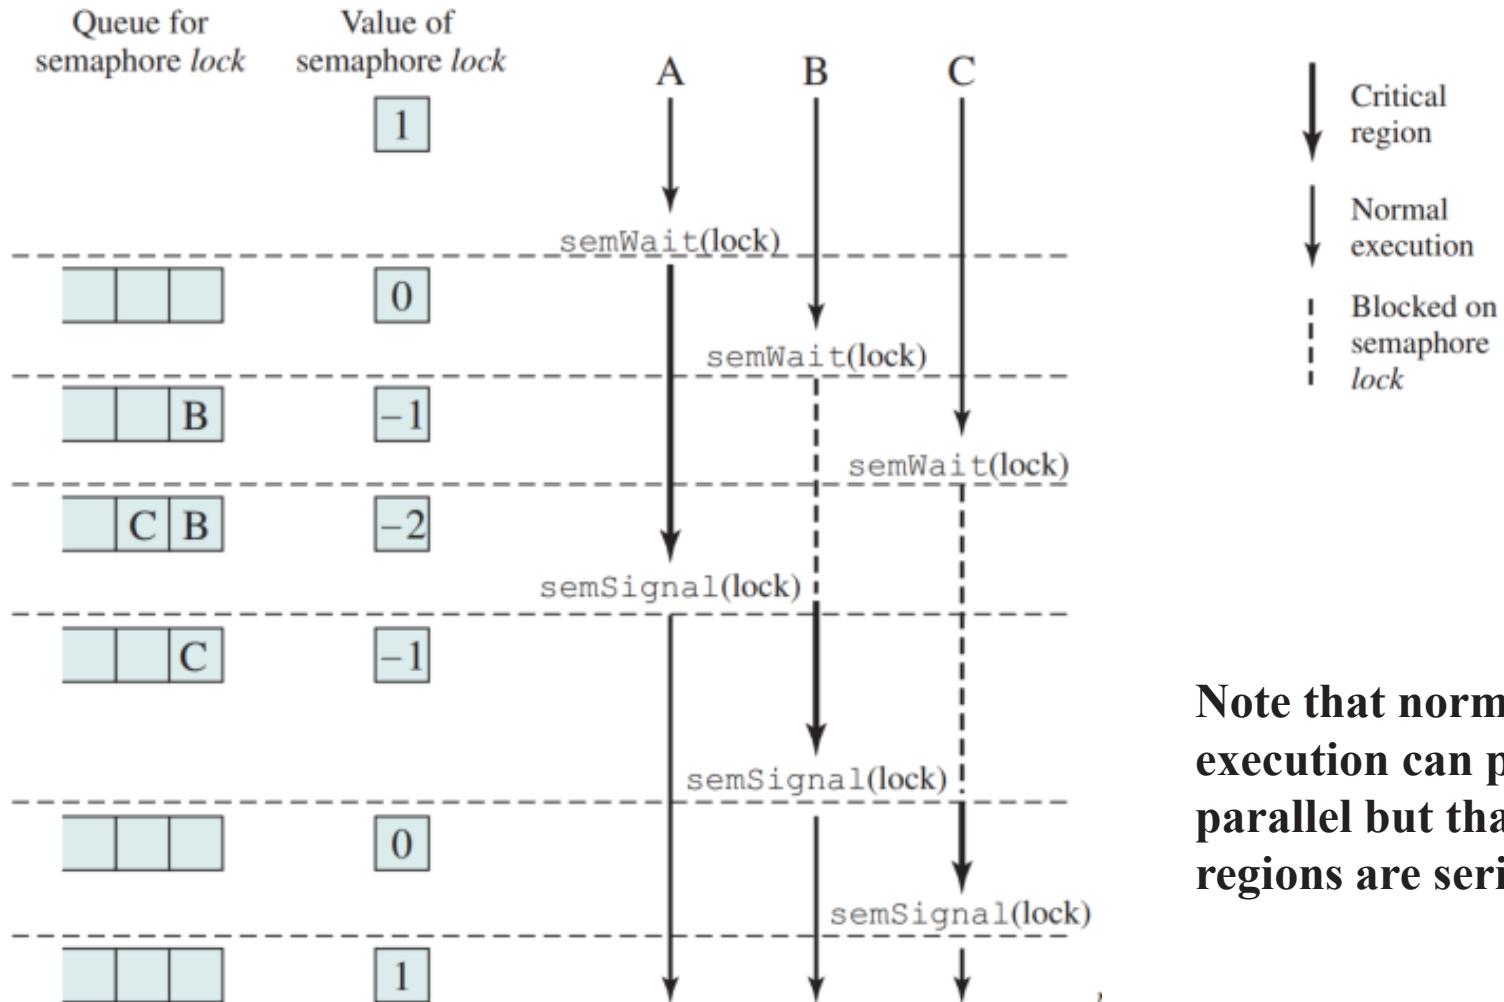


Example of Semaphore Mechanism

Mutual Exclusion Using Semaphores

```
/* program mutualexclusion */
const int n = /* number of processes */;
semaphore s = 1;
void P(int i)
{
    while (true) {
        semWait(s);
        /* critical section */;
        semSignal(s);
        /* remainder */;
    }
}
void main()
{
    parbegin (P(1), P(2), . . . , P(n));
}
```

Shared Data Protected by a Semaphore



Note that normal execution can proceed in parallel but that critical regions are serialized.

Producer/Consumer Problem

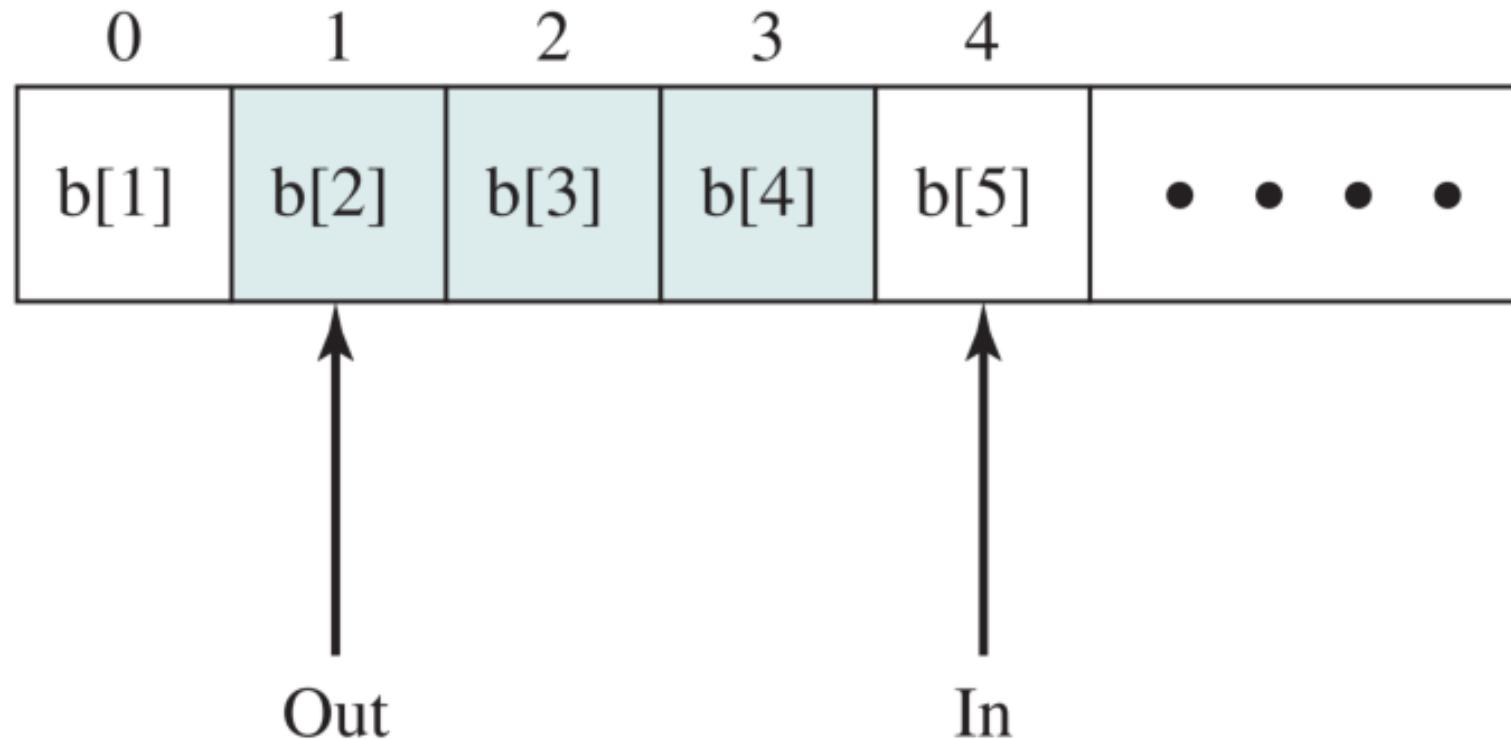
■ General Situation:

- one or more producers are generating data and placing these in a buffer
- a single consumer is taking items out of the buffer **one at time**
- **only one** producer or consumer **may access** the buffer **at any one time**

■ The problem

- ensure that the producer **can't add** data into **full** buffer
- ensure the consumer **can't remove** data from an **empty** buffer

Infinite Buffer



Note: Shaded area indicates portion of buffer that is occupied.

An Incorrect Solution to the Infinite-Buffer Producer/Consumer Problem Using Binary Semaphores

```
/* program producerconsumer */
int n;
binary_semaphore s = 1, delay = 0;
void producer()
{
    while (true) {
        produce();
        semWaitB(s);
        append();
        n++;
        if (n==1) semSignalB(delay);
        semSignalB(s);
    }
}
void consumer()
{
    semWaitB(delay);
    while (true) {
        semWaitB(s);
        take();
        n--;
        semSignalB(s);
        consume();
        if (n==0) semWaitB(delay);
    }
}
void main()
{
    n = 0;
    parbegin (producer, consumer);
}
```

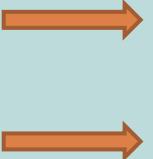
What's meaning?

Possible Scenario for the Program above

	Producer	Consumer	s	n	Delay
1			1	0	0
2	semWaitB(s)		0	0	0
3	n++		0	1	0
4	if (n==1) (semSignalB(delay))		0	1	1
5	semSignalB(s)		1	1	1
6		semWaitB(delay)	1	1	0
7		semWaitB(s)	0	1	0
8		n--	0	0	0
9		semSignalB(s)	1	0	0
10	semWaitB(s)		0	0	0
11	n++		0	1	0
12	if (n==1) (semSignalB(delay))		0	1	1
13	semSignalB(s)		1	1	1
14		if (n==0) (semWaitB(delay))	1	1	1
15		semWaitB(s)	0	1	1
16		n--	0	0	1
17		semSignalB(s)	1	0	1
18		if (n==0) (semWaitB(delay))	1	0	0
19		semWaitB(s)	0	0	0
20		n--	0	-1	0
21		semSignalB(s)	1	-1	0

A Correct Solution to the Infinite-Buffer Producer/Consumer Problem Using Binary Semaphores

```
/* program producerconsumer */
int n;
binary_semaphore s = 1, delay = 0;
void producer()
{
    while (true) {
        produce();
        semWaitB(s);
        append();
        n++;
        if (n==1) semSignalB(delay);
        semSignalB(s);
    }
}
void consumer()
{
    int m; /* a local variable */
semWaitB(delay);
while (true) {
    semWaitB(s);
    take();
    n--;
    m = n;
    semSignalB(s);
    consume();
    if (m==0) semWaitB(delay);
}
}
void main()
{
    n = 0;
    parbegin (producer, consumer);
}
```



A Solution to the Infinite-Buffer Producer/Consumer Problem Using Semaphores

```
/* program producerconsumer */
semaphore n = 0, s = 1;
void producer()
{
    while (1) {
        produce();
        semWait(s);
        append();
        semSignal(s);
        semSignal(n); } }
void consumer()
{
    while (1) {
        semWait(n);
        semWait(s);
        take();
        semSignal(s);
        consume(); } }
void main()
{
    parbegin (producer, consumer); }
```

Which one
produce a
serious flaw?



A Solution to the Bounded-Buffer Producer/Consumer Problem Using Semaphores

```
/* program boundedbuffer */
const int sizeofbuffer = /* buffer size */;
semaphore s = 1, n = 0, e = sizeofbuffer;
void producer()
{
    while (true) {
        produce();
        semWait(e);
        semWait(s);
        append();
        semSignal(s);
        semSignal(n);
    }
}
void consumer()
{
    while (true) {
        semWait(n);
        semWait(s);
        take();
        semSignal(s);
        semSignal(e);
        consume();
    }
}
void main()
{
    parbegin (producer, consumer);
}
```

What's meaning?



An Example of Semaphore Usage

例：假如系统中有**2台**打印机可用；
现有**4个**进程**P1, P2, P3, P4**都在不同时间里
以不同数量申请该设备。
S初值=2，表示共有**2台**打印机可用。
P1、P2、P3、P4为并发进程，本例假设第**1**
个被调度的为**P2**。
列出使用**P、V**操作使这**4个**进程互斥工作过程。

并发进程Pi工作序列 ()数字表示运行顺序				并发执行 顺序	当前处于 运行态 进程	所执行的 操作 (P/V)	信号量S 的值 (初值=2)	被唤醒的 进程	信号量S的 等待队列
P1	P2	P3	P4						
·	·	·	·	(1)					
·	·	·	·						
·	·	·	·						
(2) P(S)	(1) P(S)	(3) P(S)	(5) P(S)						
·	·	·	·						
打印	打印	打印	打印						
·	·	·	·						
(6) V(S)	(4) P(S)	(7) V(S)	(10) V(S)						
·	·	·	·						
·	·	·	·						
·	·	·	·						
(9) V(S)	(8) P(S)	·	·						
·	·	·	·						
打印	打印	打印	打印						
·	·	·	·						
(11) V(S)	(12) V(S)	·	·						
·	·	·	·						
·	·	·	·						
(12)									

Implementation of Semaphores

- It is imperative that the **semWait** and **semSignal** operations be implemented as **atomic primitives**
- Can be implemented in **hardware** or **firmware**
- **Software schemes** such as Dekker' s or Peterson' s algorithms can be used
- Use one of the **hardware-supported schemes** for mutual exclusion

Disadvantages of Semaphores

- **Semaphores provide a primitive yet powerful and flexible tool**

- for enforcing mutual exclusion
- for coordinating processes

- **Difficult to produce a correct program using semaphores**

- **semWait** and **semSignal** operations may be scattered throughout a program
- not easy to see the overall effect of these operations on the semaphores they affect

```
/* program producerconsumer */
int n;
binary_semaphore s = 1, delay = 0;
void producer()
{
    while (true) {
        produce();
        semWaitB(s);
        append();
        n++;
        if (n==1) semSignalB(delay);
        semSignalB(s);
    }
}
void consumer()
{
    semWaitB(delay);
    while (true) {
        semWaitB(s);
        take();
        n--;
        semSignalB(s);
        consume();
        if (n==0) semWaitB(delay);
    }
}
void main()
{
    n = 0;
    parbegin (producer, consumer);
}
```

Monitors

- **The monitor is a programming-language construct**

- that provides **equivalent** functionality to that of **semaphores**

- **First formally defined by Hoare in 1974**



For his **fundamental contributions** to the definition and design of programming languages

C. Antony ("Tony") R. Hoare
United Kingdom – 1980

- **It is easier to control**

- Implemented in many programming languages
 - including Concurrent Pascal, Modula-2, Modula-3, and Java
 - Implemented as a **program library**

- **It is a software module**

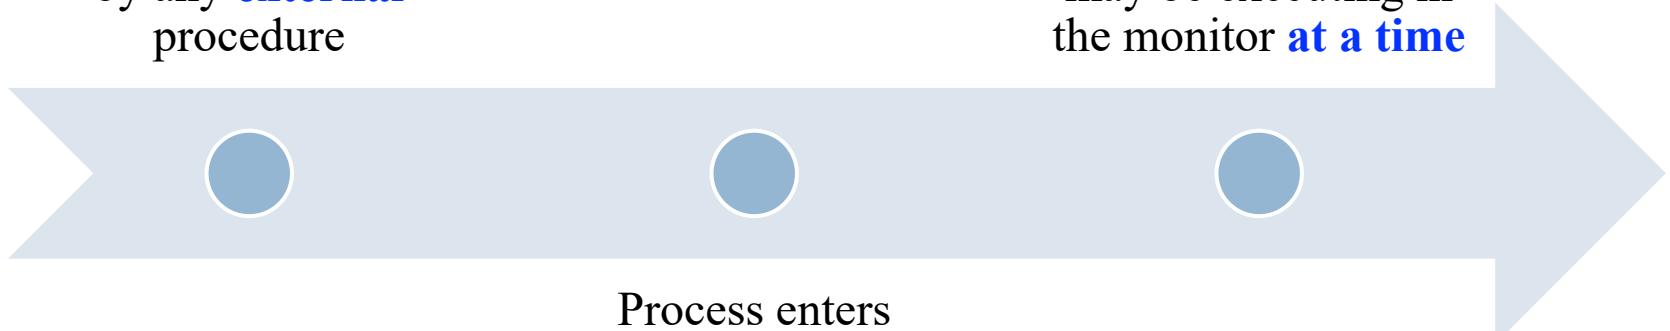
- one or more **procedures**
 - an **initialization sequence**
 - **local data**



Monitor Characteristics

Local data variables are accessible **only** by the monitor's procedures and **not** by any **external** procedure

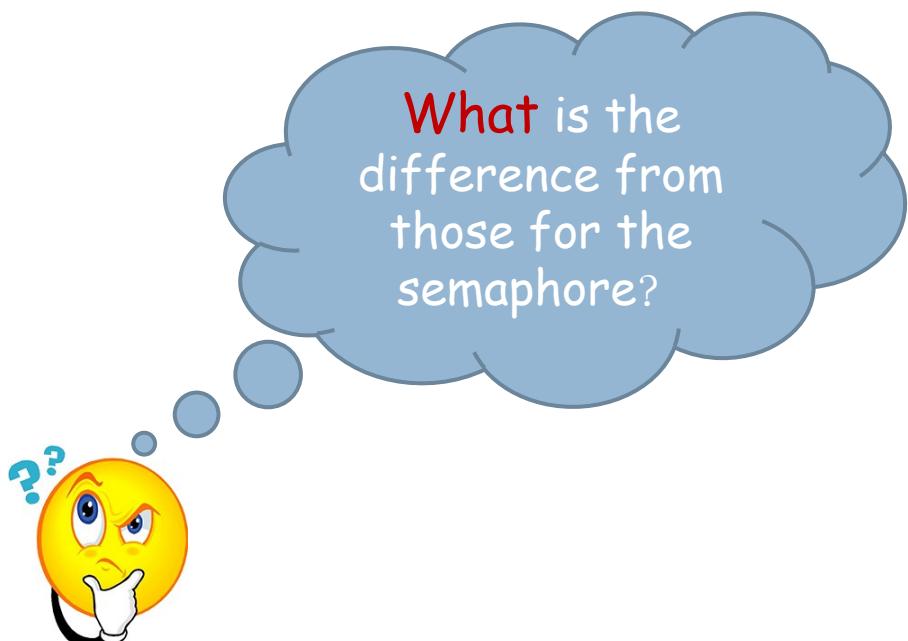
Only one process may be executing in the monitor **at a time**



Process enters monitor by **invoking one of its procedures**

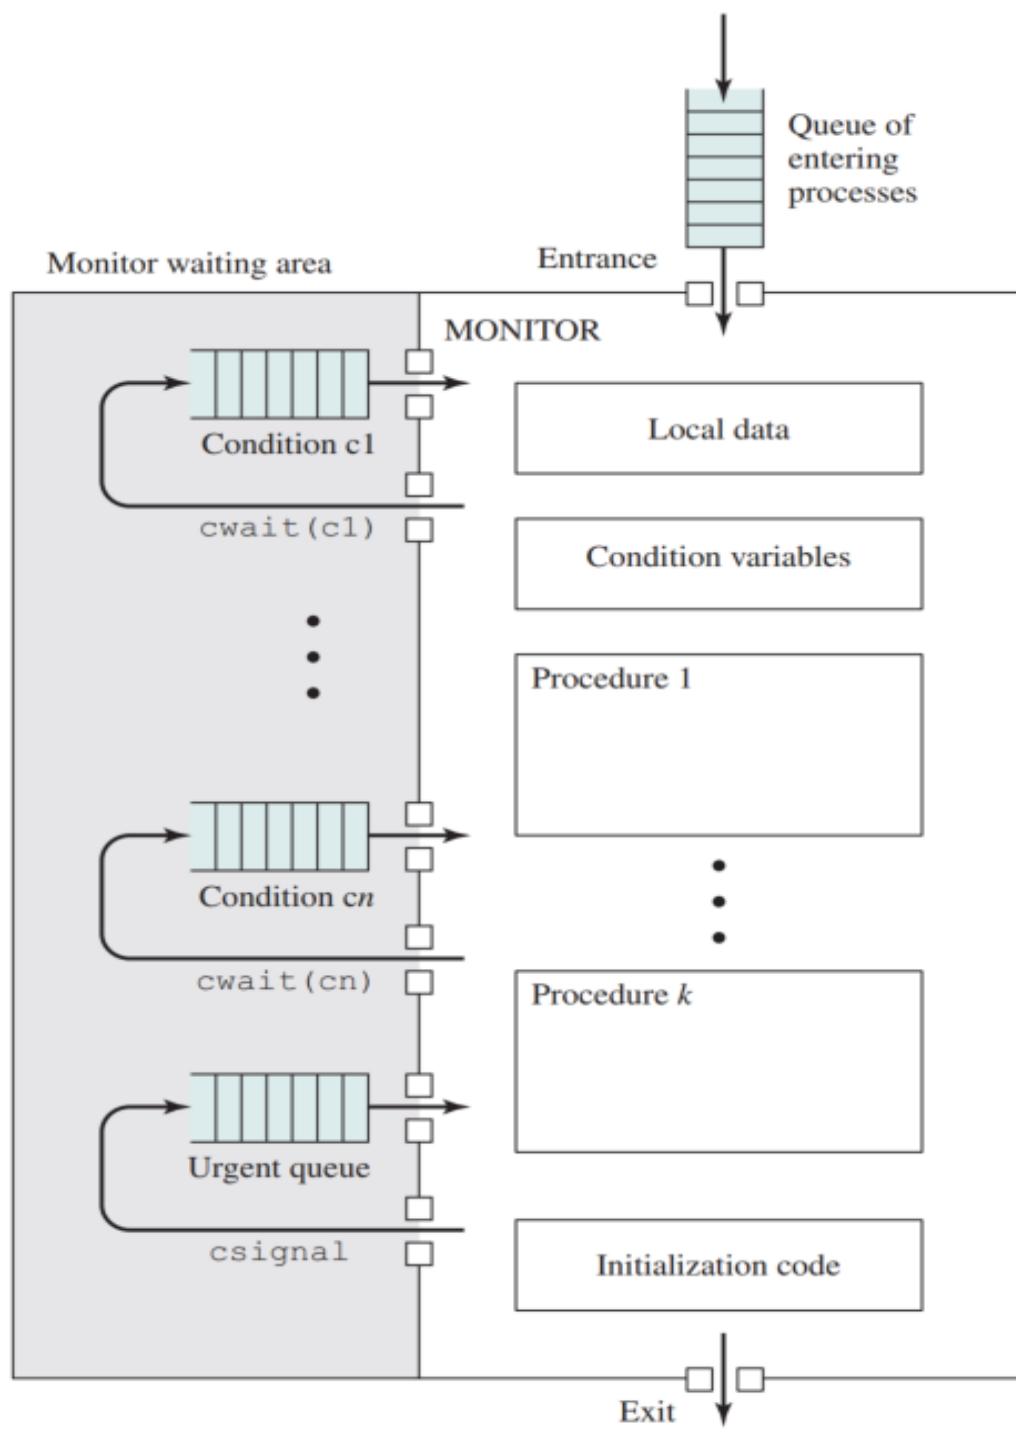
Synchronization

- Achieved by the use of **condition variables** that are contained within the monitor and **accessible only within the monitor**



- Condition variables are operated on by two functions
 - **cwait(c)**
 - suspend execution of the calling process on condition c
 - **csignal(c)**
 - resume execution of some process blocked after a cwait on the same condition
 - if there is no such process, do nothing

Structure of a Monitor



```

/* program producerconsumer */
monitor boundedbuffer;
char buffer [N];
int nextin, nextout;
int count;
cond notfull, notempty;
void append (char x)
{
    if (count == N) cwait(notfull);           /* space for N items */
    buffer[nextin] = x;                      /* buffer pointers */
    nextin = (nextin + 1) % N;                /* number of items in buffer */
    count++;                                 /* condition variables for synchronization */
    /* one more item in buffer */
    csignal (notempty);                     /* resume any waiting consumer */
}
void take (char x) —————— Procedure
{
    if (count == 0) cwait(notempty);           /* buffer is full; avoid overflow */
    x = buffer[nextout];                    /* resume any waiting consumer */
    nextout = (nextout + 1) % N;              /* one fewer item in buffer */
    count--;                                /* resume any waiting producer */
    csignal (notfull);                      /* monitor body */
}
{
    nextin = 0; nextout = 0; count = 0;      /* buffer initially empty */
}

```

Local data

Condition variable

Procedure

Initialization code

A Solution to the Bounded-Buffer Producer/Consumer Problem Using a Monitor

```
void producer()
{
    char x;
    while (true) {
        produce(x);
        append(x);
    }
}
void consumer()
{
    char x;
    while (true) {
        take(x);
        consume(x);
    }
}
void main()
{
    parbegin (producer, consumer);
}
```

A Solution to the Bounded-Buffer Producer/Consumer Problem Using a Monitor

Three advantages over Semaphore

- all of the synchronization functions are confined to the monitor
- it is easier to verify that the synchronization has been done correctly and to detect bugs
- once a monitor is correctly programmed, access to the protected resource is correct for access from all processes

Monitors with Notify and Broadcast

□ Two drawbacks to Hoare' s approach

- If the process issuing the csignal has not finished with the monitor, then **two additional process switches** are required
 - one to block this process
 - another to resume it when the monitor becomes available
- Process scheduling associated with a signal must be **perfectly reliable**



Butler W Lampson
United States – 1992

Message Passing

□ When processes interact with one another **two fundamental requirements must be satisfied**

synchronization

- to enforce mutual exclusion

communication

- to exchange information

□ **Message Passing is one approach to providing both of these functions**

- works with distributed systems and shared memory multiprocessor and uniprocessor systems

Message Passing

- The actual function is normally provided in the form of a pair of **primitives**
 - **send** (**destination**, **message**)
 - **receive** (**source**, **message**)
- A process **sends** information in the form of a message to another process designated by a destination
- A process **receives** information by executing the receive primitive, indicating the source and the message



Design Characteristics of Message Systems

□ Synchronization

- Send
 - blocking
 - nonblocking
- Receive
 - blocking
 - nonblocking
 - test for arrival

□ Format

- Content
 - Length
 - fixed
 - variable

□ Addressing

- Direct
 - send
 - receive
 - explicit
 - implicit
- Indirect
 - static
 - dynamic
 - ownership

□ Queueing Discipline

- FIFO
- Priority

Synchronization

The communication of a message between two processes implies some level of **synchronization** between the two

The receiver cannot receive a message until it has been sent by another process

When a **receive primitive** is executed in a process, there are two possibilities

If a message has previously been sent, the message is received and execution continues.

If there is no waiting message, then either (a) the process is blocked until a message arrives, or (b) the process continues to execute, abandoning the attempt to receive.

Blocking Send, Blocking Receive

❑ Both sender and receiver are blocked until the message is delivered

- Sometimes referred to as a **rendezvous**
- Allows for **tight synchronization** between processes

Nonblocking Send

Nonblocking send, blocking receive

- sender continues on but receiver is blocked until the requested message arrives
- **most useful combination**
- sends one or more messages to a variety of destinations as quickly as possible
- example -- a **service process** that exists to provide a service or resource to other processes

Nonblocking send, nonblocking receive

- neither party is required to wait

Direct Addressing

- Send primitive includes a specific identifier of the destination process
- Receive primitive can be handled in one of two ways:
 - require that the process explicitly designate a sending process
 - effective for cooperating concurrent processes
 - implicit addressing
 - source parameter of the receive primitive possesses a value returned when the receive operation has been performed



Indirect Addressing

Messages are sent to a shared data structure consisting of queues that can temporarily hold messages

Queues are referred to as **mailboxes**

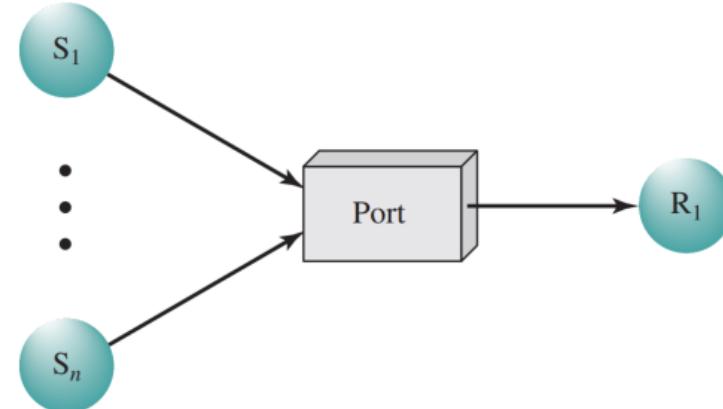
Allows for greater flexibility in the use of messages

One process sends a message to the mailbox and the other process picks up the message from the mailbox

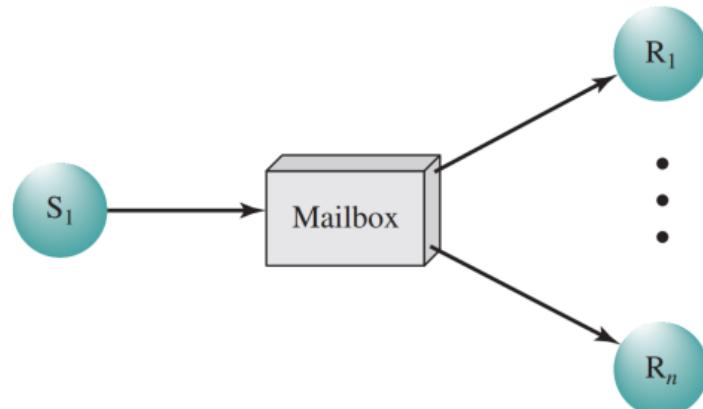
Indirect Process Communication



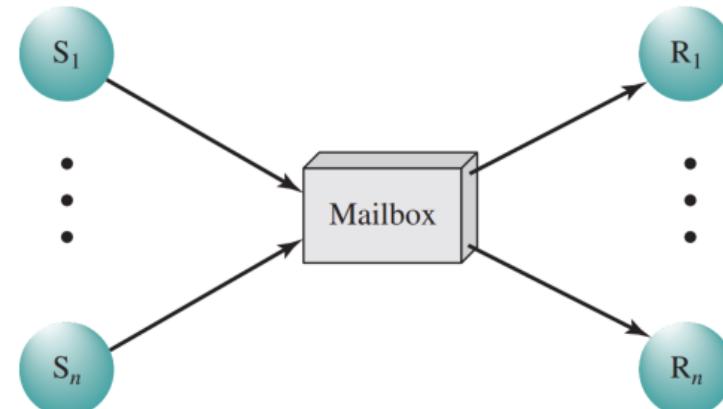
(a) One-to-one



(b) Many-to-one

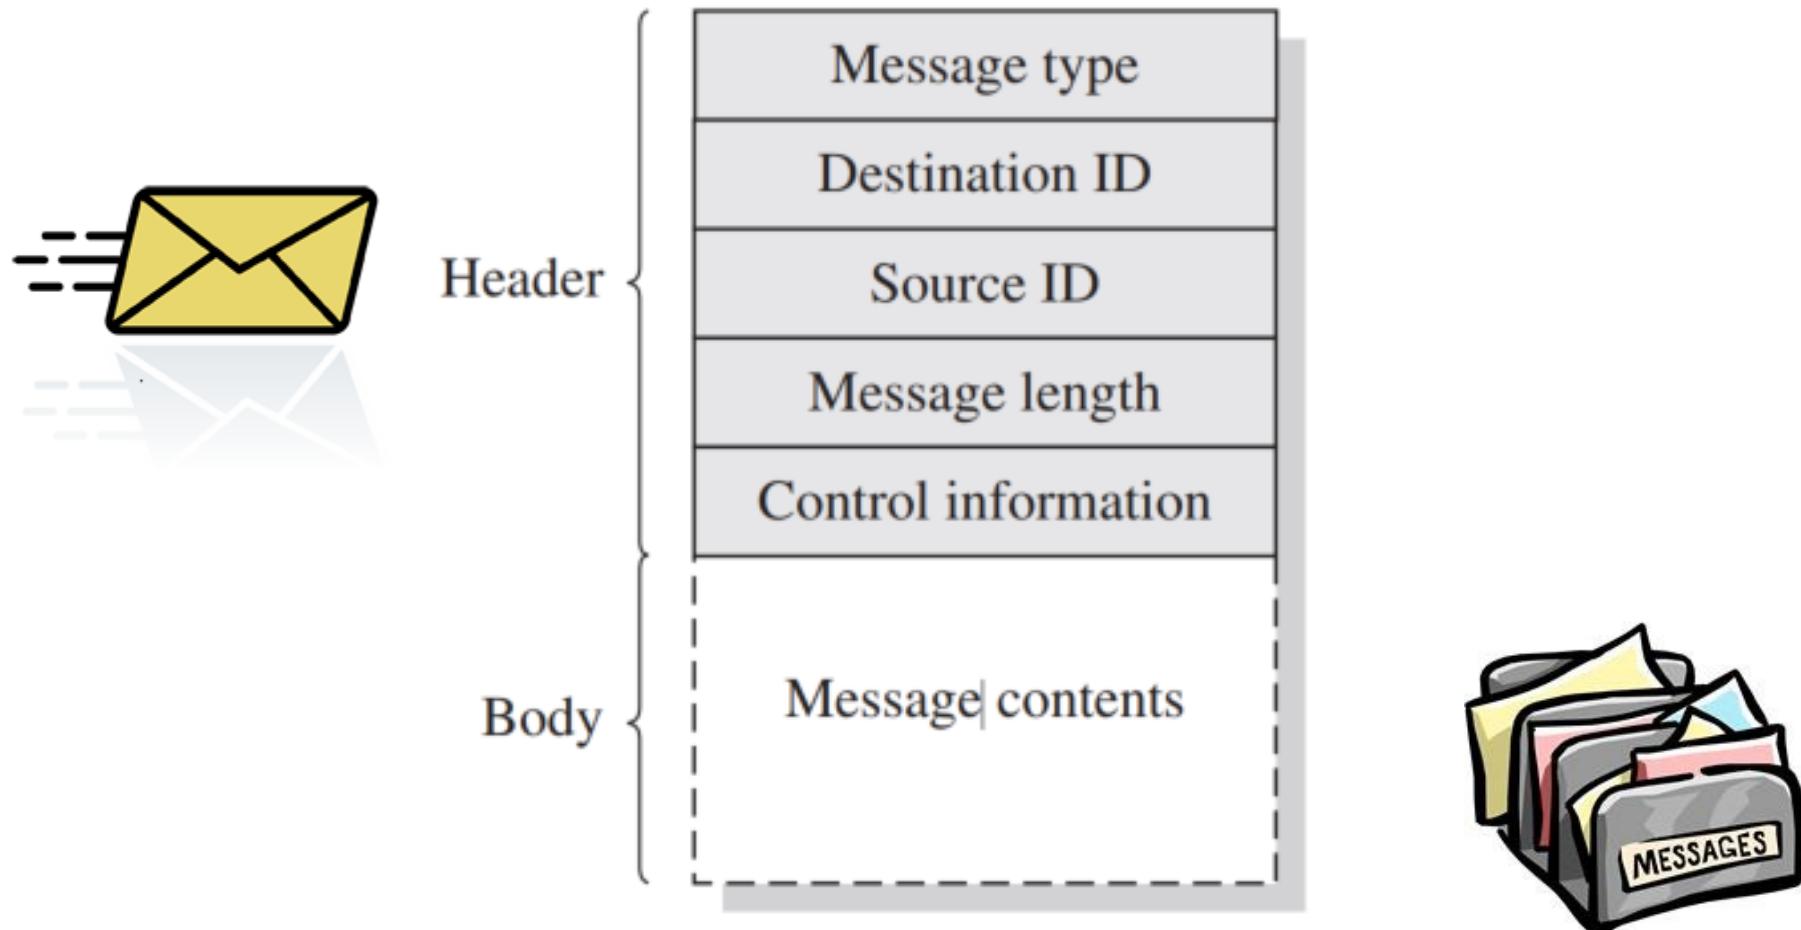


(c) One-to-many



(d) Many-to-many

General Message Format



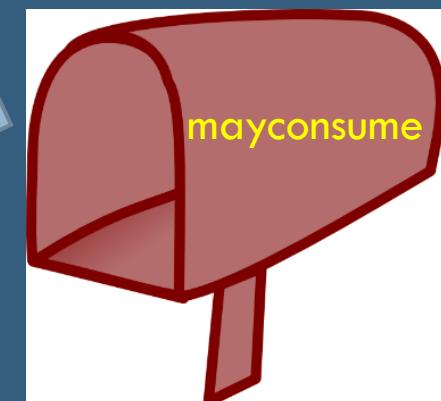
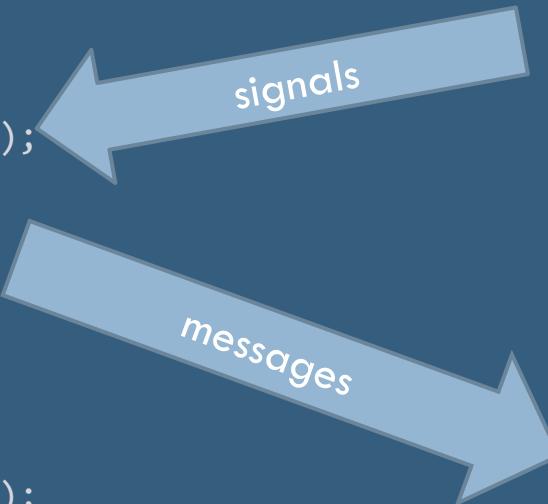
Mutual Exclusion Using Messages

```
/* program mutualExclusion */
const int n = /* number of process */
void P(int i)
{
    message msg;
    while (true) {
        receive (box, msg);
        /* critical section */;
        send (box, msg);
        /* remainder */;
    }
}
void main()
{
    create mailbox (box);
    send (box, null);
    parbegin (P(1), P(2), . . . , P(n));
}
```

A Solution to the Bounded-Buffer Producer/Consumer Problem Using a Messages

```
const int capacity = /* buffering capacity */ ;
null = /* empty message */ ;
int i;
void producer()
{   message pmsg;
    while (true) {
        receive (mayproduce,pmsg);
        pmsg = produce();
        send (mayconsume,pmsg);
    }
}
void consumer()
{   message cmsg;
    while (true) {
        receive (mayconsume,cmsg);
        consume (cmsg);
        send (mayproduce,null);
    }
}
void main()
{
    create_mailbox (mayproduce);
    create_mailbox (mayconsume);
    for (int i = 1;i<= capacity;i++) send (mayproduce,null);
    parbegin (producer,consumer);
}
```

What's meaning?



Readers/Writers Problem

❑ A data area is **shared** among many processes

- some processes **only read** the data area (readers)
- some **only write** to the data area (writers)

❑ Conditions that must be satisfied:

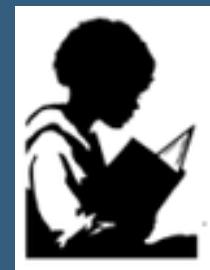
- any number of readers may **simultaneously read** the file
- **only one writer** at **a time** may write to the file
- if a writer is writing to the file, **no reader** may read it

Readers/Writers Problem



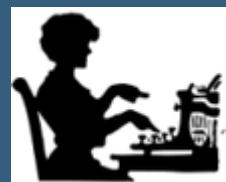
A Solution to the Readers/Writers Problem Using Semaphores: Readers Have Priority

```
/* program readersandwriters */
int readcount;
semaphore x = 1, wsem = 1;
void reader()
{
    while (true){
        semWait(x);
        readcount++;
        if (readcount == 1) semWait(wsem);
        semSignal(x);
        READUNIT();
        semWait(x);
        readcount--;
        if (readcount == 0) semSignal(wsem);
        semSignal(x);
    }
}
void writer()
{
    while (true){
        semWait(wsem);
        WRITEUNIT();
        semSignal(wsem);
    }
}
void main()
{
    readcount = 0;
    parbegin(reader, writer);
}
```



A Solution to the Readers/Writers Problem Using Semaphores: Writer Have Priority

```
int readcount, writecount;
semaphore x = 1, y = 1, z = 1, wsem = 1, rsem = 1;
void reader()
{
    while (true){
        semWait(z);
        semWait(rsem);
        semWait(x);
        readcount++;
        if (readcount == 1) semWait(wsem);
        semSignal(x);
        semSignal(rsem);
        semSignal(z);
        READUNIT();
        semWait(x);
        readcount--;
        if (readcount == 0) semSignal(wsem);
        semSignal(x);
    }
}
void writer()
{
    while (true){
        semWait(y);
        writecount++;
        if (writecount == 1) semWait(rsem);
        semSignal(y);
        semWait(wsem);
        WRITEUNIT();
        semSignal(wsem);
        semWait(y);
        writecount--;
        if (writecount == 0) semSignal(rsem);
        semSignal(y);
    }
}
void main()
{
    readcount = writecount = 0;
    parbegin(reader, writer);
}
```



State of the Process Queues

Readers only in the system

- wsem set
- no queues

Writers only in the system

- wsem and rsem set
- writers queue on wsem

Both readers and writers with read first

- wsem set by reader
- rsem set by writer
- all writers queue on wsem
- one reader queues on rsem
- other readers queue on z

Both readers and writers with write first

- wsem set by writer
- rsem set by writer
- writers queue on wsem
- one reader queues on rsem
- other readers queue on z

A Solution to the Readers/Writers Problem Using Message Passing

```
void reader(int i)                                void controller()
{
    message rmsg;
    while (true) {
        rmsg = i;
        send (readrequest, rmsg);
        receive (mbox[i], rmsg);
        READUNIT ();
        rmsg = i;
        send (finished, rmsg);
    }
}

void writer(int j)
{
    message rmsg;
    while (true){
        rmsg = j;
        send (writerequest, rmsg);
        receive (mbox[j], rmsg);
        WRITEUNIT ();
        rmsg = j;
        send (finished, rmsg);
    }
}

void controller()
{
    while (true)
    {
        if (count > 0) {
            if (!empty (finished)) {
                receive (finished, msg);
                count++;
            }
        } else if (!empty (writerequest)) {
            receive (writerequest, msg);
            writer_id = msg.id;
            count = count - 100;
        } else if (!empty (readrequest)) {
            receive (readrequest, msg);
            count--;
            send (msg.id, "OK");
        }
        if (count == 0) {
            send (writer_id, "OK");
            receive (finished, msg);
            count = 100;
        }
        while (count < 0) {
            receive (finished, msg);
            count++;
        }
    }
}
```

Summary

- **Operating system themes**
 - Multiprogramming, multiprocessing, distributed processing
 - **Fundamental** to these themes is **concurrency**
 - issues of **conflict** resolution and **cooperation** arise
- **Mutual Exclusion**
 - Condition in which there is a set of **concurrent processes, only one** of which is able to access a given resource or perform a given function at any time
 - Three approaches to supporting
- **Mutual Exclusion: software approaches**
- **Hardware support**
- **OS or a programming language**
 - **Semaphores**
 - Used for **signaling** among processes and can be readily used to **enforce a mutual exclusion discipline**
 - **Monitors**
 - **Messages**
 - Useful for the **enforcement of mutual exclusion discipline**
 - provide an effective means of **interprocess communication**