

操作系统

第6章 并发性：死锁和饥饿 Concurrency: Deadlock and Starvation

孙承杰
哈工大计算学部

E-mail: sunchengjie@hit.edu.cn
2025年秋季学期

Learning Objectives

- List and explain the **conditions for deadlock**
- Define deadlock prevention and describe deadlock prevention strategies related to each of the conditions for deadlock
- Explain the **difference** between deadlock prevention and deadlock avoidance
- Understand two approaches to deadlock avoidance
- Explain the **fundamental difference** in approach between deadlock detection and deadlock prevention or avoidance
- Analyze the **dining philosophers problem**

Outline

■ Principles of Deadlock

- Reusable resources
- Consumable resources
- Resource allocation graphs
- The conditions for deadlock

■ Deadlock Prevention

- Mutual exclusion
- Hold and wait
- No preemption
- Circular wait

■ Deadlock Avoidance

- Process initiation denial
- Resource allocation denial

■ Deadlock Detection

- Deadlock detection algorithm
- Recovery

■ Dining Philosophers Problem

- Solution using semaphores
- Solution using a monitor

When **two trains** approach each other at a crossing,
both shall come to a full stop and neither shall start up
again until the other has gone. Statute passed by the
Kansas State Legislature, early in the 20th century.

-- *A TREASURY OF RAILROAD FOLKLORE,*
B. A. Botkin and Alvin F. Harlow

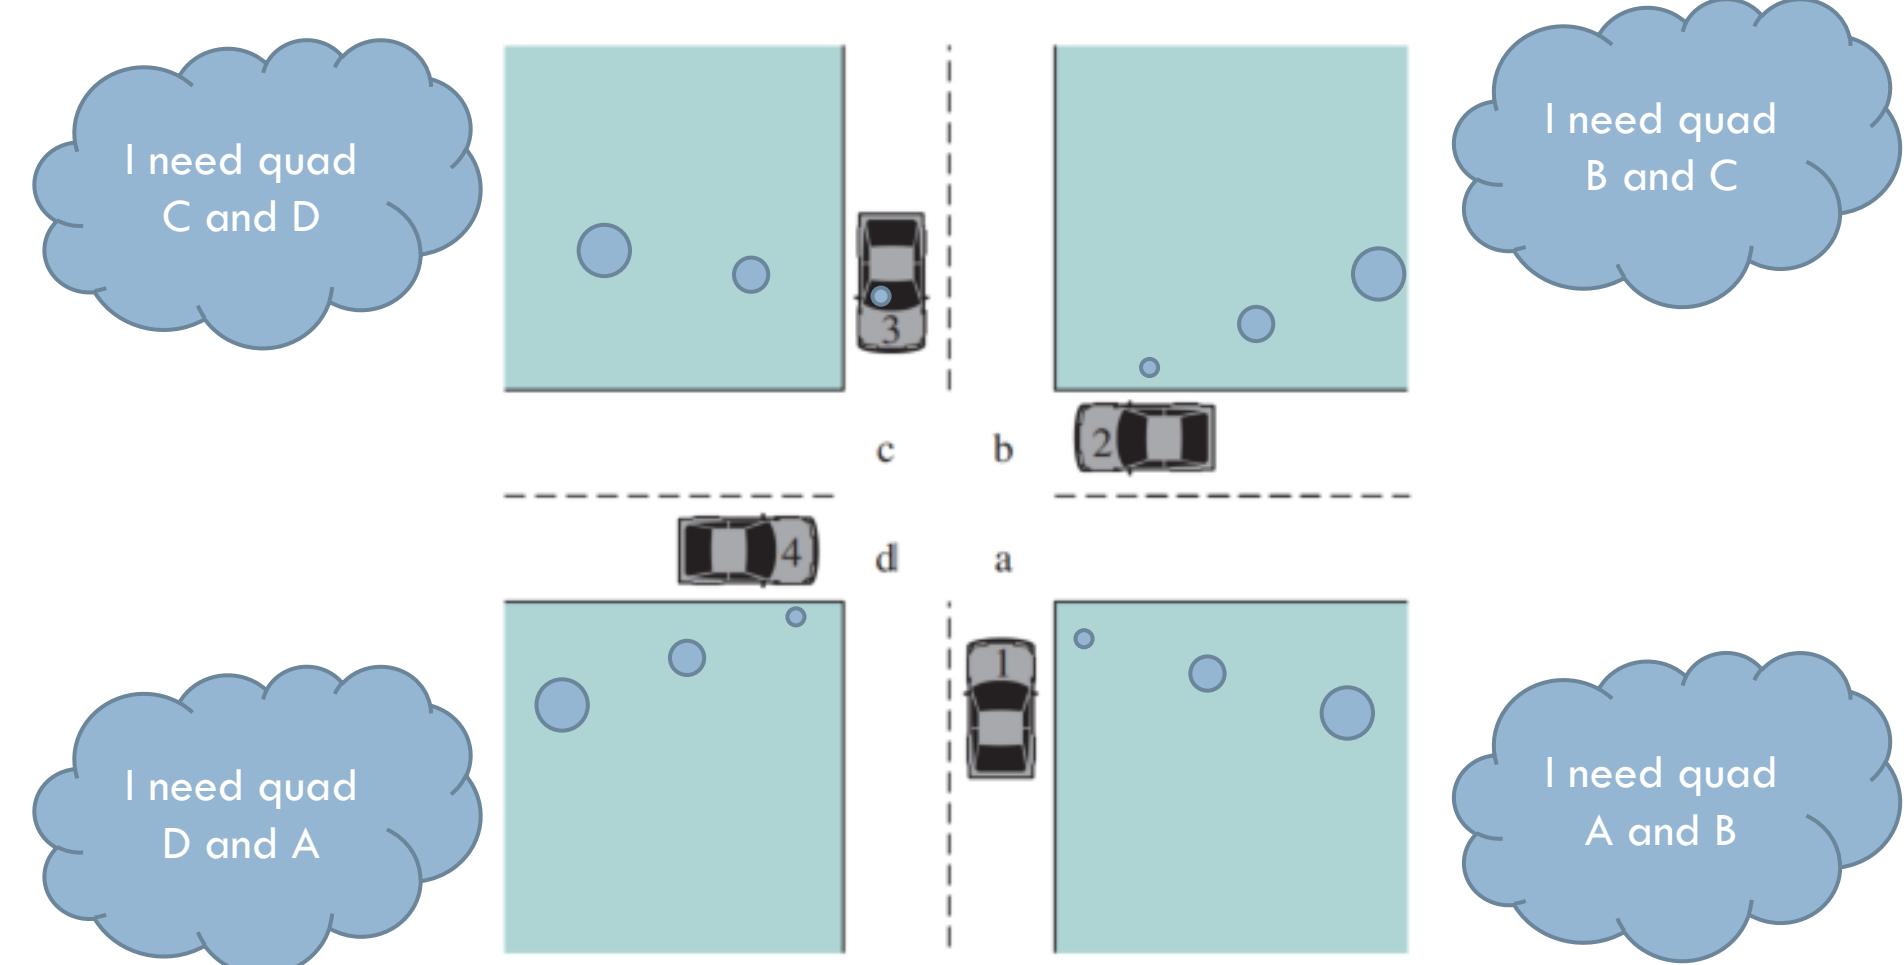


Deadlock

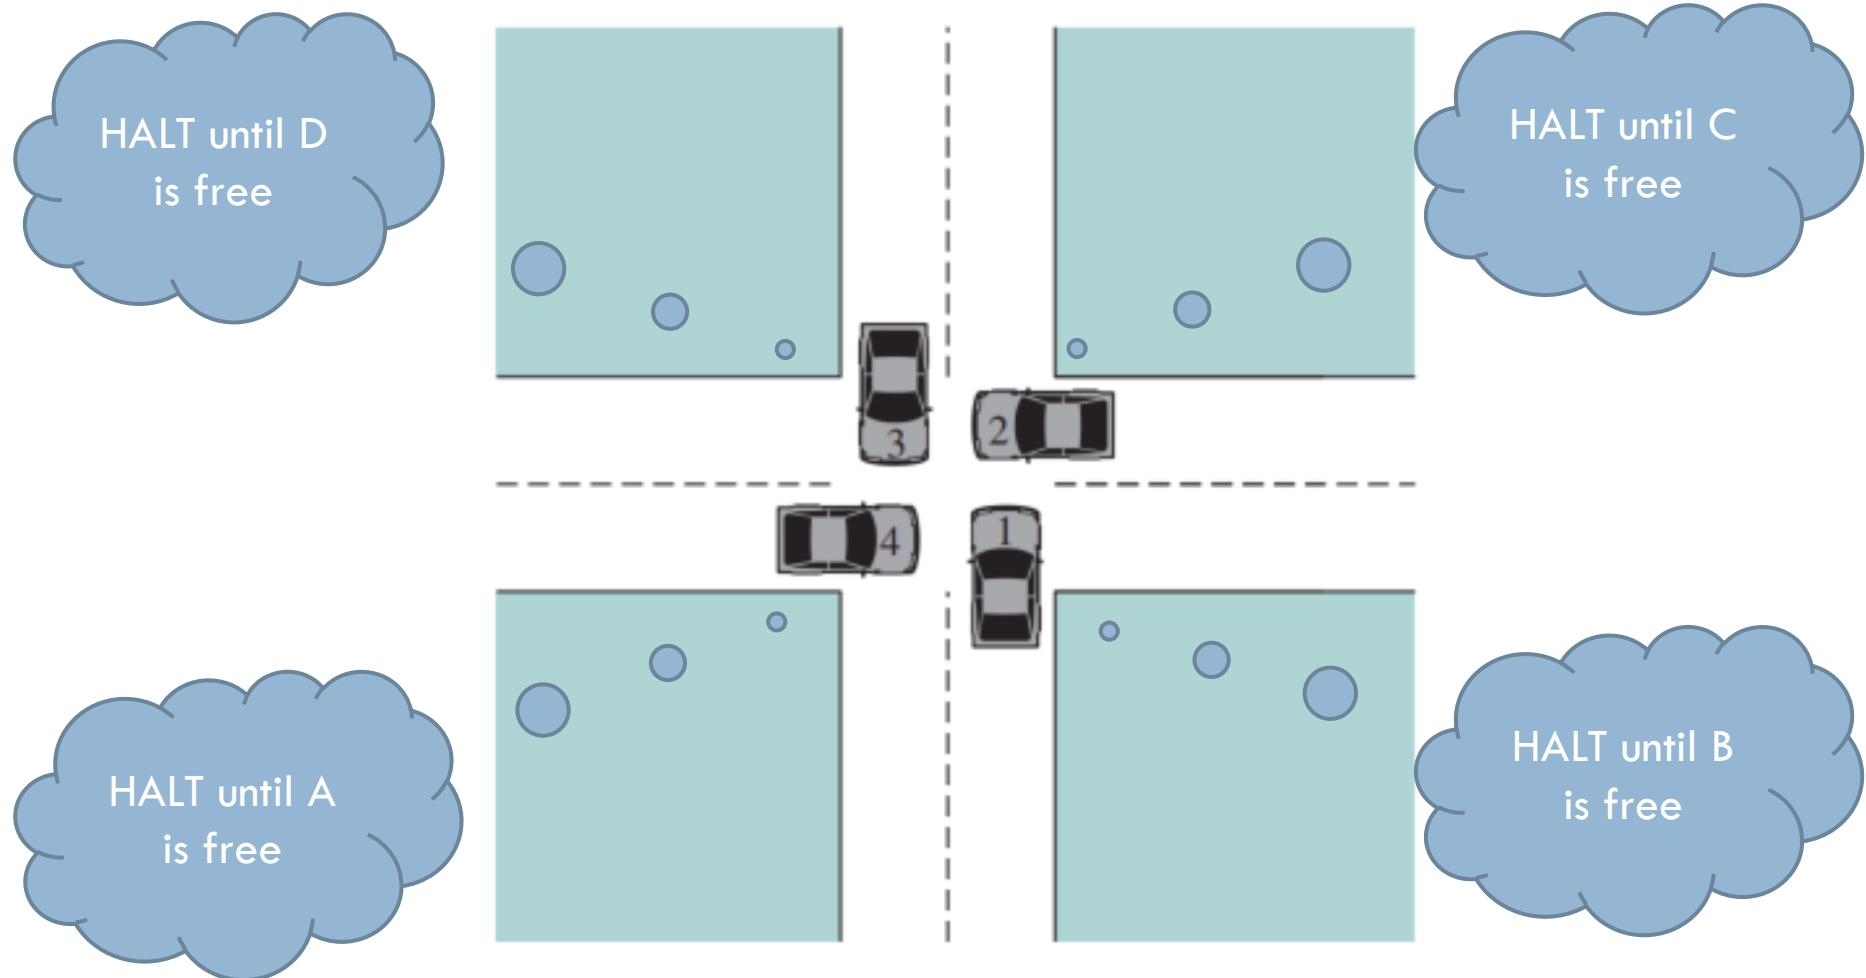
- The **permanent blocking** of a set of processes that either compete for system resources or communicate with each other
- A set of processes is **deadlocked** when each process in the set is blocked awaiting an event that can only be triggered by another blocked process in the set
- Permanent
- No efficient solution



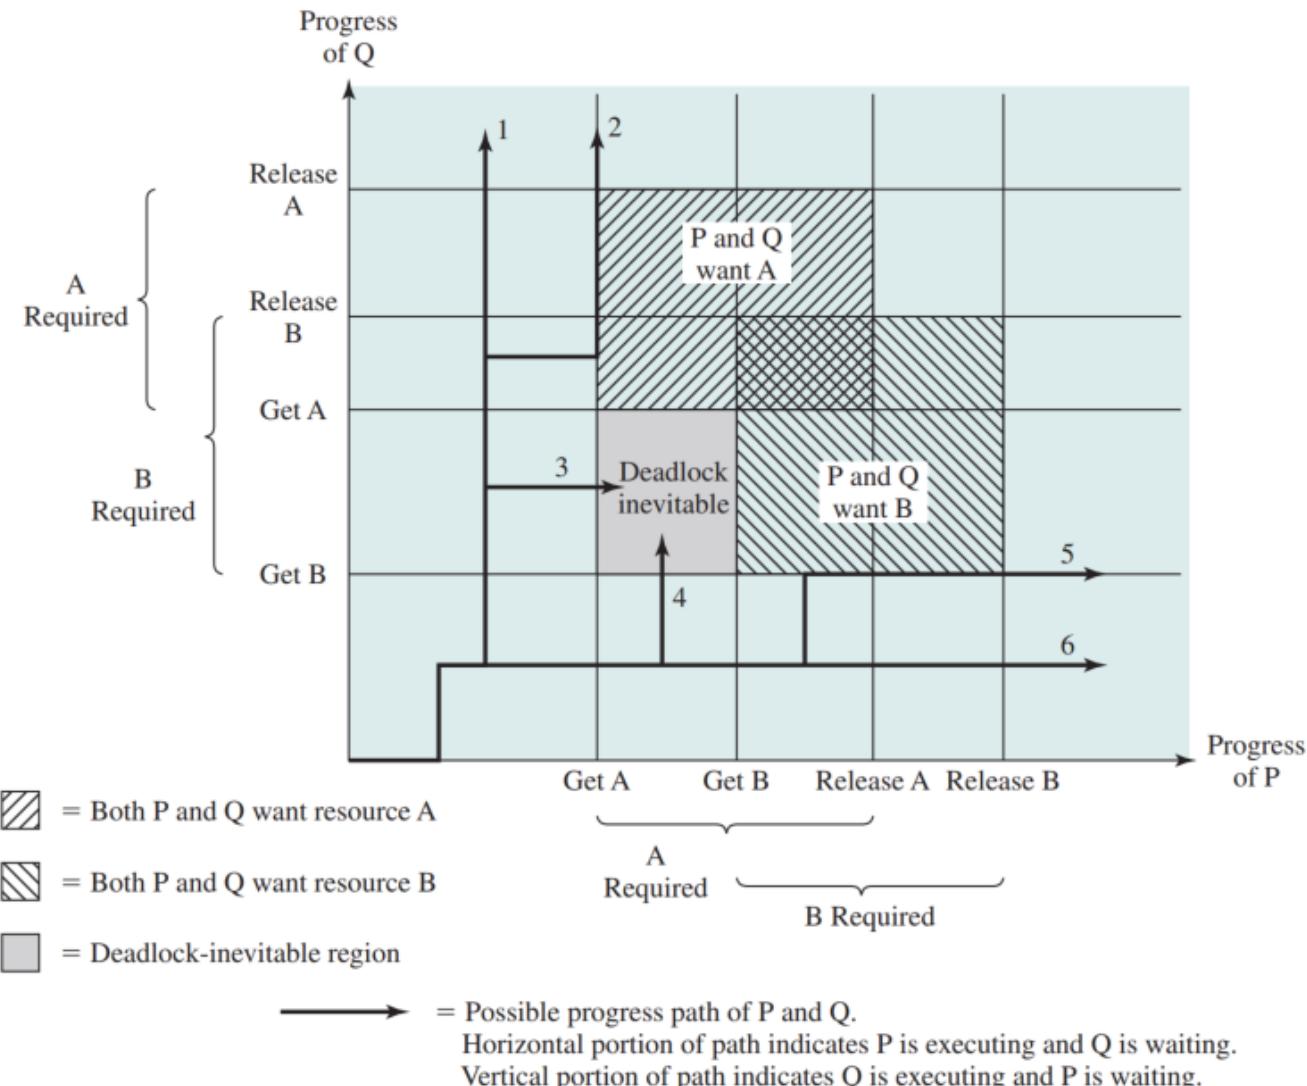
Potential Deadlock



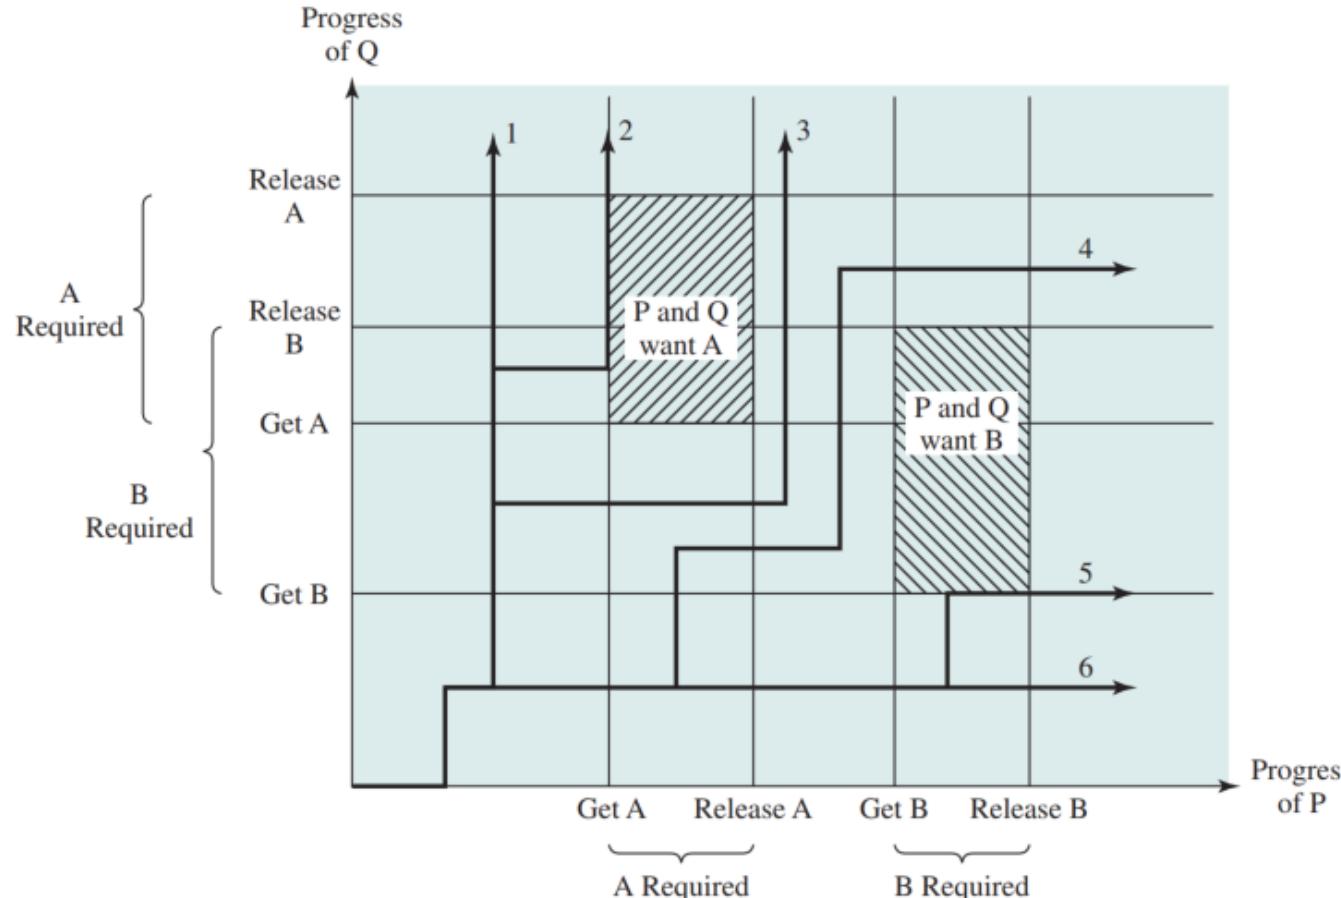
Actual Deadlock



Deadlock Example



No Deadlock Example



■ = Both P and Q want resource A

■ = Both P and Q want resource B → = Possible progress path of P and Q.
Horizontal portion of path indicates P is executing and Q is waiting.
Vertical portion of path indicates Q is executing and P is waiting.

Resource Categories

■ Reusable(可重用)

- can be safely used by **only one process at a time** and is **not depleted** by that use
 - processors
 - I/O channels
 - main and secondary memory
 - devices
 - data structures such as files, databases, and semaphores

■ Consumable(可消耗)

- one that can be **created** (produced) and **destroyed** (consumed)
 - interrupts
 - signals
 - messages
 - information in I/O buffers

Reusable Resources: Example 1

Step	Process P Action
p ₀	Request (D)
p ₁	Lock (D)
p ₂	Request (T)
p ₃	Lock (T)
p ₄	Perform function
p ₅	Unlock (D)
p ₆	Unlock (T)

Step	Process Q Action
q ₀	Request (T)
q ₁	Lock (T)
q ₂	Request (D)
q ₃	Lock (D)
q ₄	Perform function
q ₅	Unlock (T)
q ₆	Unlock (D)

Example 2: Memory Request

- Space is available for allocation of **200 Kbytes**, and the following sequence of events occur

P1	P2
...	...
Request 80 Kbytes;	Request 70 Kbytes;
...	...
Request 60 Kbytes;	Request 80 Kbytes;

- Deadlock occurs if both processes progress to their second request

Consumable Resources Deadlock

- Consider a pair of processes, in which each process attempts to receive a message from the other process and then send a message to the other process:

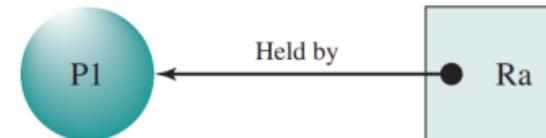
P1	P2
...	...
Receive (P2);	Receive (P1);
...	...
Send (P2, M1);	Send (P1, M2);

- Deadlock occurs if the **Receive is blocking**

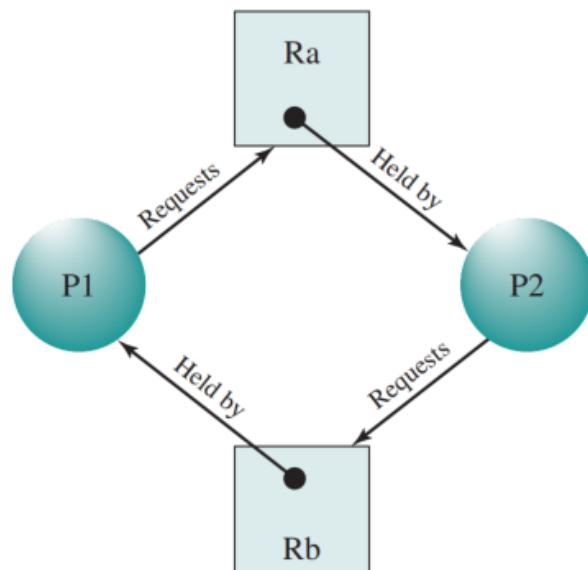
Resource Allocation Graphs



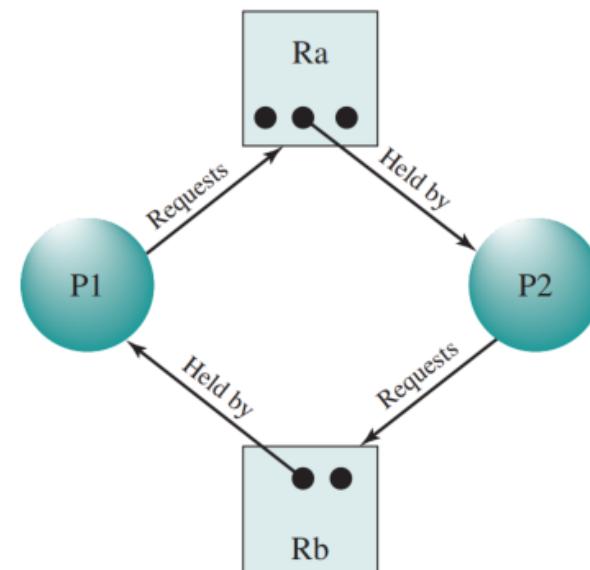
(a) Resource is requested



(b) Resource is held

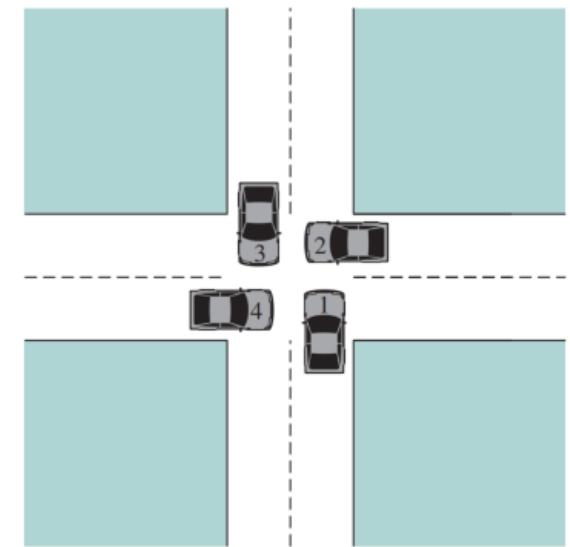
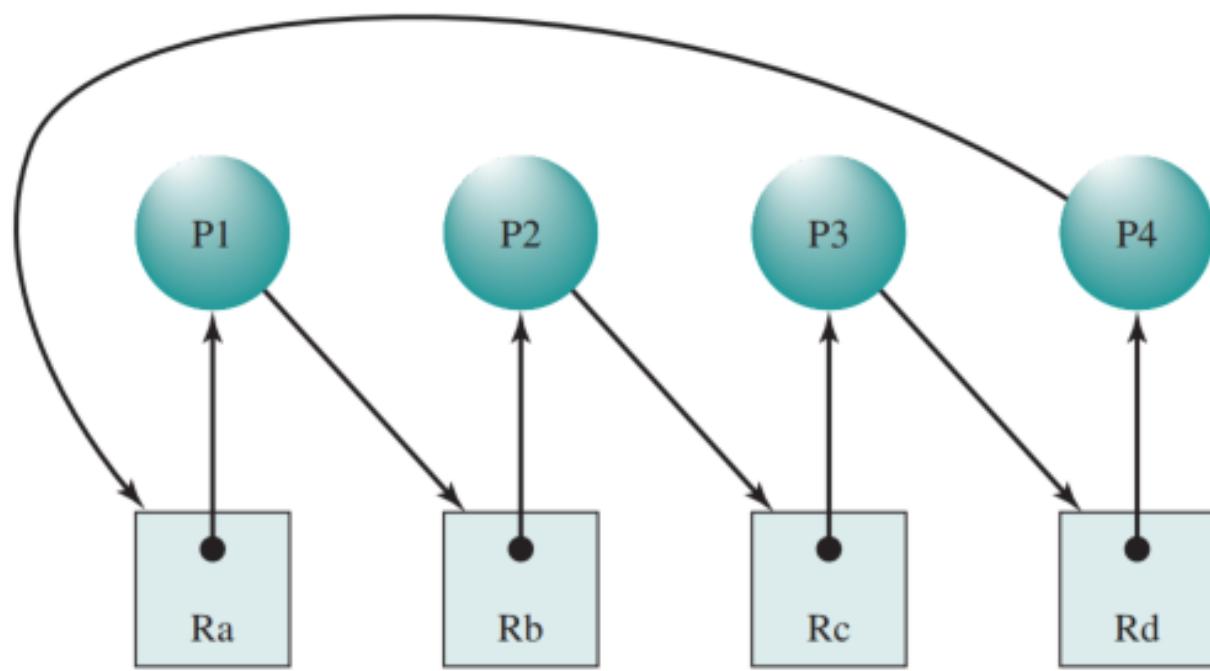


(c) Circular wait



(d) No deadlock

Resource Allocation Graphs



Conditions for Deadlock

Mutual exclusion

- Only one process may use a resource at a time. No process may access a resource unit that has been allocated to another process.

Policy decisions

Hold and wait

- A process may hold allocated resources while awaiting assignment of other resources.

No preemption

- No resource can be forcibly removed from a process holding it.

Circular wait

- A closed chain of processes exists, such that each process holds at least one resource needed by the next process in the chain.

Might occur

To summarize

- **Possibility of Deadlock**
 - Mutual exclusion
 - No preemption
 - Hold and wait
- **Existence of Deadlock**
 - Mutual exclusion
 - No preemption
 - Hold and wait
 - **Circular wait**



Dealing with Deadlock

Prevent Deadlock(死锁预防)

- by adopting a policy that **eliminates one of the conditions**

Avoid Deadlock(死锁避免)

- make the **appropriate dynamic choices** based on the current state of resource allocation

Detect Deadlock(死锁检测)

- attempt to **detect** the presence of deadlock and take action to recover

Deadlock Prevention Strategy

□ Design a system in such a way that the possibility of deadlock is excluded

□ Two main methods

- Indirect
 - prevent the occurrence of **one of the three** necessary conditions
- Direct
 - prevent the occurrence of a **circular wait**

Deadlock Condition Prevention

■ Mutual Exclusion

- if access to a resource requires mutual exclusion then it must be supported by the OS

■ Hold and Wait

- require that a process request all of its required resources **at one time** and blocking the process until all requests can be **granted simultaneously**

■ No Preemption

- if a process holding certain resources is denied a further request, that process must **release** its original resources and request them again
- OS may preempt the second process and require it to release its resources
 - **different priority**

■ Circular Wait

- define a **linear ordering** of resource types

Deadlock Avoidance

- A decision is made **dynamically**

- whether the current resource allocation request will, if granted, potentially lead to a deadlock

- Requires knowledge of future process requests

Different from deadlock prevention allows the three necessary conditions but makes judicious(明智的) choices to assure that the **deadlock point is never reached**



Two Approaches to Deadlock Avoidance

Deadlock Avoidance

**Process Initiation
Denial:**
Do not start a process if its demands might lead to deadlock

**Resource Allocation
Denial:**
Do not grant an incremental resource request to a process if this allocation might lead to deadlock

Resource Allocation Denial

□ Referred to as **the banker's algorithm**

□ **State of the system**

- reflects the current allocation of resources to processes

□ **Safe state**

- is one in which there is **at least one sequence** of resource allocations to processes that **does not result in a deadlock**

□ **Unsafe state**

- is a state that is not safe

死锁避免之银行家算法

一个银行家：目前手里只有1亿，但是已经贷出很多钱

开发商A：已贷款15亿，资金紧张还需3亿。

开发商B：已贷款5亿，还需贷款1亿，运转良好能收回。

开发商C：已贷款2亿，欲贷款18亿

会不会出现楼盘烂尾？

开发商B还钱，再借给A，则可以继续借给C

银行家当前手里现金（Available）；

银行家可以利用的资金，即手里现金加上能收回的共有多少（work）；

各个开发商已贷款——已分配的资金（Allocation）；

各个开发商还需要贷款（Claim- Allocation）

钱就是资源，开发商就是进程，银行家的决策就是调度

A system of n processes and m resources

Resource = $\mathbf{R} = (R_1, R_2, \dots, R_m)$	Total amount of each resource in the system
Available = $\mathbf{V} = (V_1, V_2, \dots, V_m)$	Total amount of each resource not allocated to any process
Claim = $\mathbf{C} = \begin{pmatrix} C_{11} & C_{12} & \dots & C_{1m} \\ C_{21} & C_{22} & \dots & C_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ C_{n1} & C_{n2} & \dots & C_{nm} \end{pmatrix}$	C_{ij} = requirement of process i for resource j
Allocation = $\mathbf{A} = \begin{pmatrix} A_{11} & A_{12} & \dots & A_{1m} \\ A_{21} & A_{22} & \dots & A_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ A_{n1} & A_{n2} & \dots & A_{nm} \end{pmatrix}$	A_{ij} = current allocation to process i of resource j

1. $R_j = V_j + \sum_{i=1}^n A_{ij}$, for all j All resources are either available or allocated.

2. $C_{ij} \leq R_j$, for all i, j No process can claim more than the total amount of resources in the system.

3. $A_{ij} \leq C_{ij}$, for all i, j No process is allocated more resources of any type than the process originally claimed to need.

Determination of a Safe State

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	1	0	0
P2	6	1	2
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	2	2	2
P2	0	0	1
P3	1	0	3
P4	4	2	0

$C - A$

	R1	R2	R3
	9	3	6

Resource vector R

	R1	R2	R3
	0	1	1

Available vector V

(a) Initial state

	R1	R2	R3
P1	3	2	2
P2	0	0	0
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	1	0	0
P2	0	0	0
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	2	2	2
P2	0	0	0
P3	1	0	3
P4	4	2	0

$C - A$

	R1	R2	R3
	9	3	6

Resource vector R

	R1	R2	R3
	6	2	3

Available vector V

(b) P2 runs to completion

Determination of a Safe State

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	1	0	3
P4	4	2	0

C - A

	R1	R2	R3
	9	3	6

Resource vector R

	R1	R2	R3
	7	2	3

Available vector V

(c) P1 runs to completion

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	4	2	0

C - A

	R1	R2	R3
	9	3	6

Resource vector R

	R1	R2	R3
	9	3	4

Available vector V

(d) P3 runs to completion

Determination of a Unsafe State

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	1	0	0
P2	5	1	1
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	2	2	2
P2	1	0	2
P3	1	0	3
P4	4	2	0

$C - A$

	R1	R2	R3
	9	3	6

Resource vector R

	R1	R2	R3
	1	1	2

Available vector V

(a) Initial state

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim matrix C

	R1	R2	R3
P1	2	0	1
P2	5	1	1
P3	2	1	1
P4	0	0	2

Allocation matrix A

	R1	R2	R3
P1	1	2	1
P2	1	0	2
P3	1	0	3
P4	4	2	0

$C - A$

	R1	R2	R3
	9	3	6

Resource vector R

	R1	R2	R3
	0	1	1

Available vector V

(b) P1 requests one unit each of R1 and R3

Deadlock Avoidance Logic

```
struct state {  
    int resource[m];  
    int available[m];  
    int claim[n][m];  
    int alloc[n][m];  
}
```

Global data structures

```
if (alloc [i,*] + request [*] > claim [i,*])  
    <error>; /* total request > claim */  
else if (request [*] > available [*])  
    <suspend process>;  
else {  
    /* simulate alloc */  
    <define newstate by:  
    alloc [i,*] = alloc [i,*] + request [*];  
    available [*] = available [*] - request [*]>;  
}  
if (safe (newstate))  
    <carry out allocation>;  
else {  
    <restore original state>;  
    <suspend process>;  
}
```

Deadlock Avoidance Logic

```
boolean safe (state S) {
    int currentavail[m];
    process rest[<number of processes>];
    currentavail = available;
    rest = {all processes};
    possible = true;
    while (possible) {
        <find a process Pk in rest such that
        claim [k,*] - alloc [k,*]<= currentavail;
        if (found) {           /* simulate execution of Pk */
            currentavail = currentavail + alloc [k,*];
            rest = rest - {Pk};
        }
        else possible = false;
    }
    return (rest == null);
}
```

Test for safety algorithm (banker's algorithm)

Advantages and Restrictions

■ Advantages

- It is **not necessary** to preempt and rollback processes, as in deadlock detection
- It is **less restrictive** than deadlock prevention

■ Restrictions

- Maximum resource requirement for each process must be stated in advance
- Processes under consideration must be independent and with no synchronization requirements
- There must be a **fixed number** of resources to allocate
- No process may exit while holding resources

Deadlock Strategies

■ Deadlock prevention

- very **conservative**
- **limit access** to resources by imposing restrictions on processes

■ Deadlock detection

- do the opposite
- resource requests are granted **whenever possible**

Deadline Detection Algorithms

- A check for deadlock can be made as **frequently** as each **resource request** or, less frequently, depending on how likely it is for a deadlock to occur
- **Advantages:**
 - it leads to **early detection**
 - the algorithm is **relatively simple**
- **Disadvantage**
 - **frequent** checks
 - **consume** considerable processor time

Deadline Detection Algorithms

Example for Deadlock Detection

	R1	R2	R3	R4	R5
P1	0	1	0	0	1
P2	0	0	1	0	1
P3	0	0	0	0	1
P4	1	0	1	0	1

Request matrix Q

	R1	R2	R3	R4	R5
P1	1	0	1	1	0
P2	1	1	0	0	0
P3	0	0	0	1	0
P4	0	0	0	0	0

Allocation matrix A

	R1	R2	R3	R4	R5
Resource vector	2	1	1	2	1
	R1	R2	R3	R4	R5
Available vector	0	0	0	0	1

The algorithm concludes with P1 and P2 unmarked, indicating these processes are deadlocked.

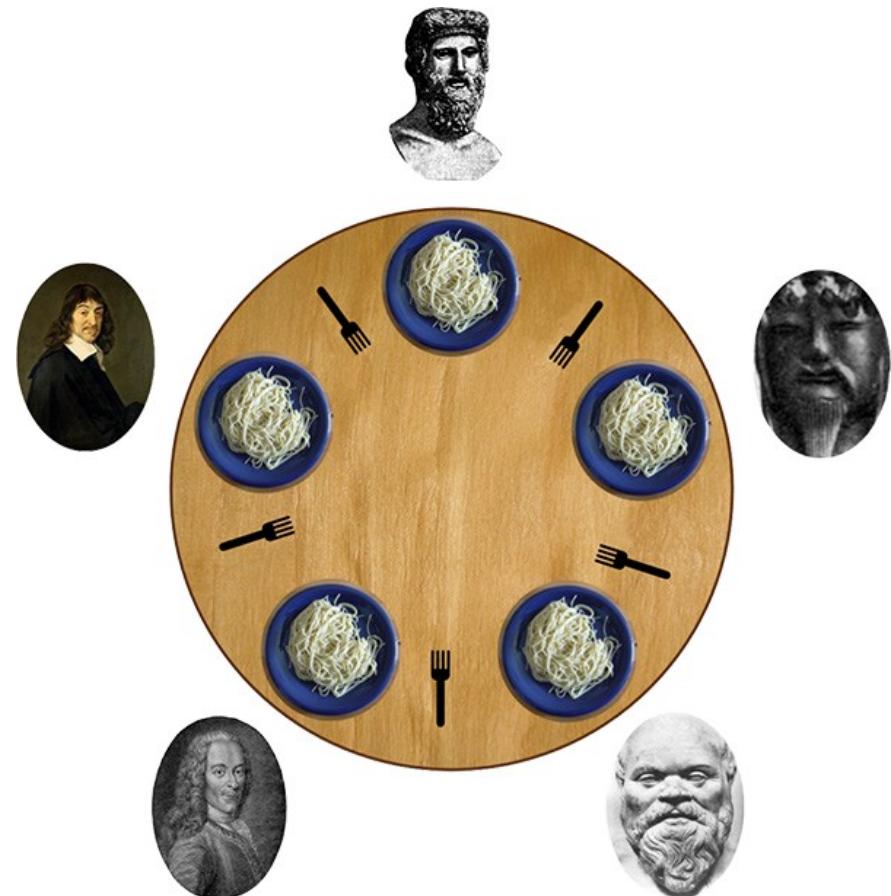
Approach	Resource Allocation Policy	Different Schemes	Major Advantages	Major Disadvantages
Prevention	Conservative; under commits resources	Requesting all resources at once	<ul style="list-style-type: none"> • Works well for processes that perform a single burst of activity • No preemption necessary 	<ul style="list-style-type: none"> • Inefficient • Delays process initiation • Future resource requirements must be known by processes
		Preemption	<ul style="list-style-type: none"> • Convenient when applied to resources whose state can be saved and restored easily 	<ul style="list-style-type: none"> • Preempts more often than necessary
		Resource ordering	<ul style="list-style-type: none"> • Feasible to enforce via compile-time checks • Needs no run-time computation since problem is solved in system design 	<ul style="list-style-type: none"> • Disallows incremental resource requests
Avoidance	Midway between that of detection and prevention	Manipulate to find at least one safe path	<ul style="list-style-type: none"> • No preemption necessary 	<ul style="list-style-type: none"> • Future resource requirements must be known by OS • Processes can be blocked for long periods
Detection	Very liberal; requested resources are granted where possible	Invoke periodically to test for deadlock	<ul style="list-style-type: none"> • Never delays process initiation • Facilitates online handling 	<ul style="list-style-type: none"> • Inherent preemption losses

Summary of Deadlock Detection, Prevention, and Avoidance Approaches for Operating Systems



Dining Philosophers Problem

- Mutual exclusion
 - No two philosophers can use the same fork at the same time
- Avoid deadlock and starvation
 - No philosopher must starve to death



Using Semaphores

```
/* program diningphilosophers */
semaphore fork [5] = {1};
int i;
void philosopher (int i)
{
    while (true) {
        think();
        wait (fork[i]);
        wait (fork [(i+1) mod 5]);
        eat();
        signal(fork [(i+1) mod 5]);
        signal(fork[i]);
    }
}
void main()
{
    parbegin (philosopher (0), philosopher (1),
              philosopher (2), philosopher (3),
              philosopher (4));
}
```

A First Solution to the Dining Philosophers Problem

Using Semaphores

A Second Solution

```
/* program diningphilosophers */
semaphore fork[5] = {1};
semaphore room = {4};
int i;
void philosopher (int i)
{
    while (true) {
        think();
        wait (room);
        wait (fork[i]);
        wait (fork [(i+1) mod 5]);
        eat();
        signal (fork [(i+1) mod 5]);
        signal (fork[i]);
        signal (room);
    }
}
void main()
{
    parbegin (philosopher (0), philosopher (1),
              philosopher (2), philosopher (3),
              philosopher (4));
}
```

Using A Monitor

```
void philosopher[k=0 to 4] /* the five philosopher clients */
{
    while (true) {
        <think>;
        get_forks(k); /* client requests two forks via monitor */
        <eat spaghetti>;
        release_forks(k); /* client releases forks via the monitor */
    }
}
```

```
monitor dining_controller;
cond ForkReady[5];      /* condition variable for synchronization */
boolean fork[5] = {true}; /* availability status of each fork */

void get_forks(int pid)    /* pid is the philosopher id number */
{
    int left = pid;
    int right = (++pid) % 5;
    /*grant the left fork*/
    if (!fork[left])
        cwait(ForkReady[left]); /* queue on condition variable */
    fork[left] = false;
    /*grant the right fork*/
    if (!fork[right])
        cwait(ForkReady[right]);/* queue on condition variable */
    fork[right] = false;
}
void release_forks(int pid)
{
    int left = pid;
    int right = (++pid) % 5;
    /*release the left fork*/
    if (empty(ForkReady[left])) /*no one is waiting for this fork */
        fork[left] = true;
    else                      /* awaken a process waiting on this fork */
        csignal(ForkReady[left]);
    /*release the right fork*/
    if (empty(ForkReady[right]))/*no one is waiting for this fork */
        fork[right] = true;
    else                      /* awaken a process waiting on this fork */
        csignal(ForkReady[right]);
}
```

Summary

■ Deadlock

- the blocking of a set of processes that either compete for system resources or communicate with each other
- blockage is permanent unless OS takes action
- may involve reusable or consumable resources
 - Consumable = destroyed
 - Reusable = not depleted or destroyed by use

■ Dealing with deadlock:

➤ prevention

- guarantees that deadlock will not occur

➤ avoidance

- analyzes each new resource request

➤ detection

- OS checks for deadlock and takes action